



Designing a Complete CI/CD Pipeline Using Argo Events, Workflows, and CD

Julian Mazzitelli, CTO BioBox Analytics Inc.

\$ whoami



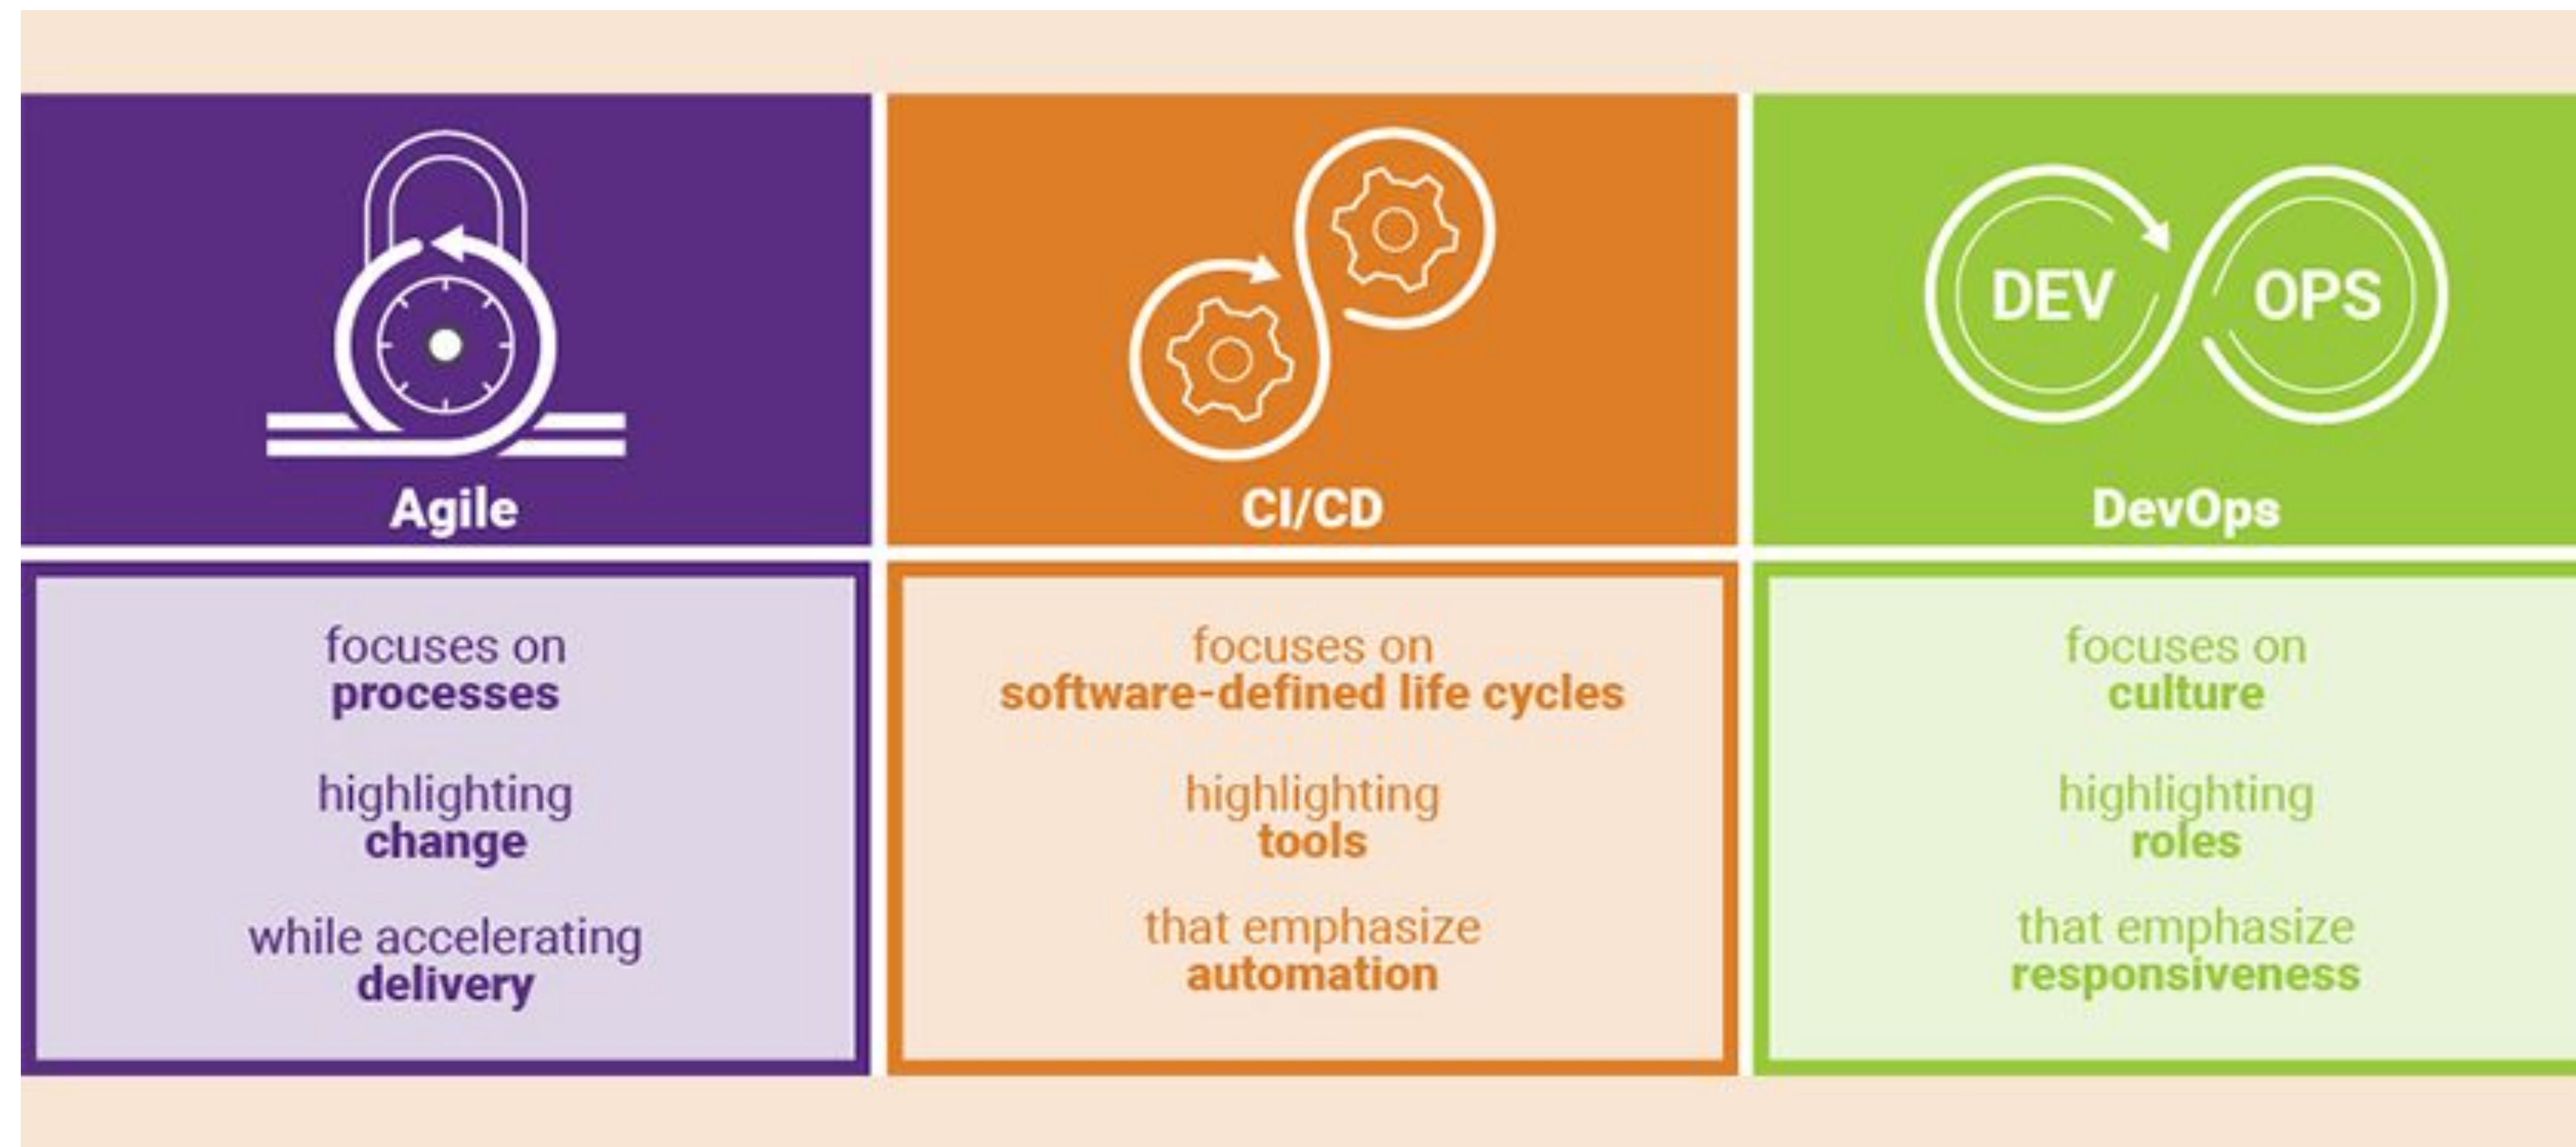
- BioBox Analytics Inc.
- Early stage startup, incorporated April 2019
- 3 full time
 - 2x developers
 - 1x developer + operations “full stack”
- 2 part time
 - Quality Assurance / Product Officer
- Cloud native stack - API talks to K8s
- WE NEED TO MOVE FAST!
 - Want a robust and flexible CI/CD process
 - Want Kubernetes native




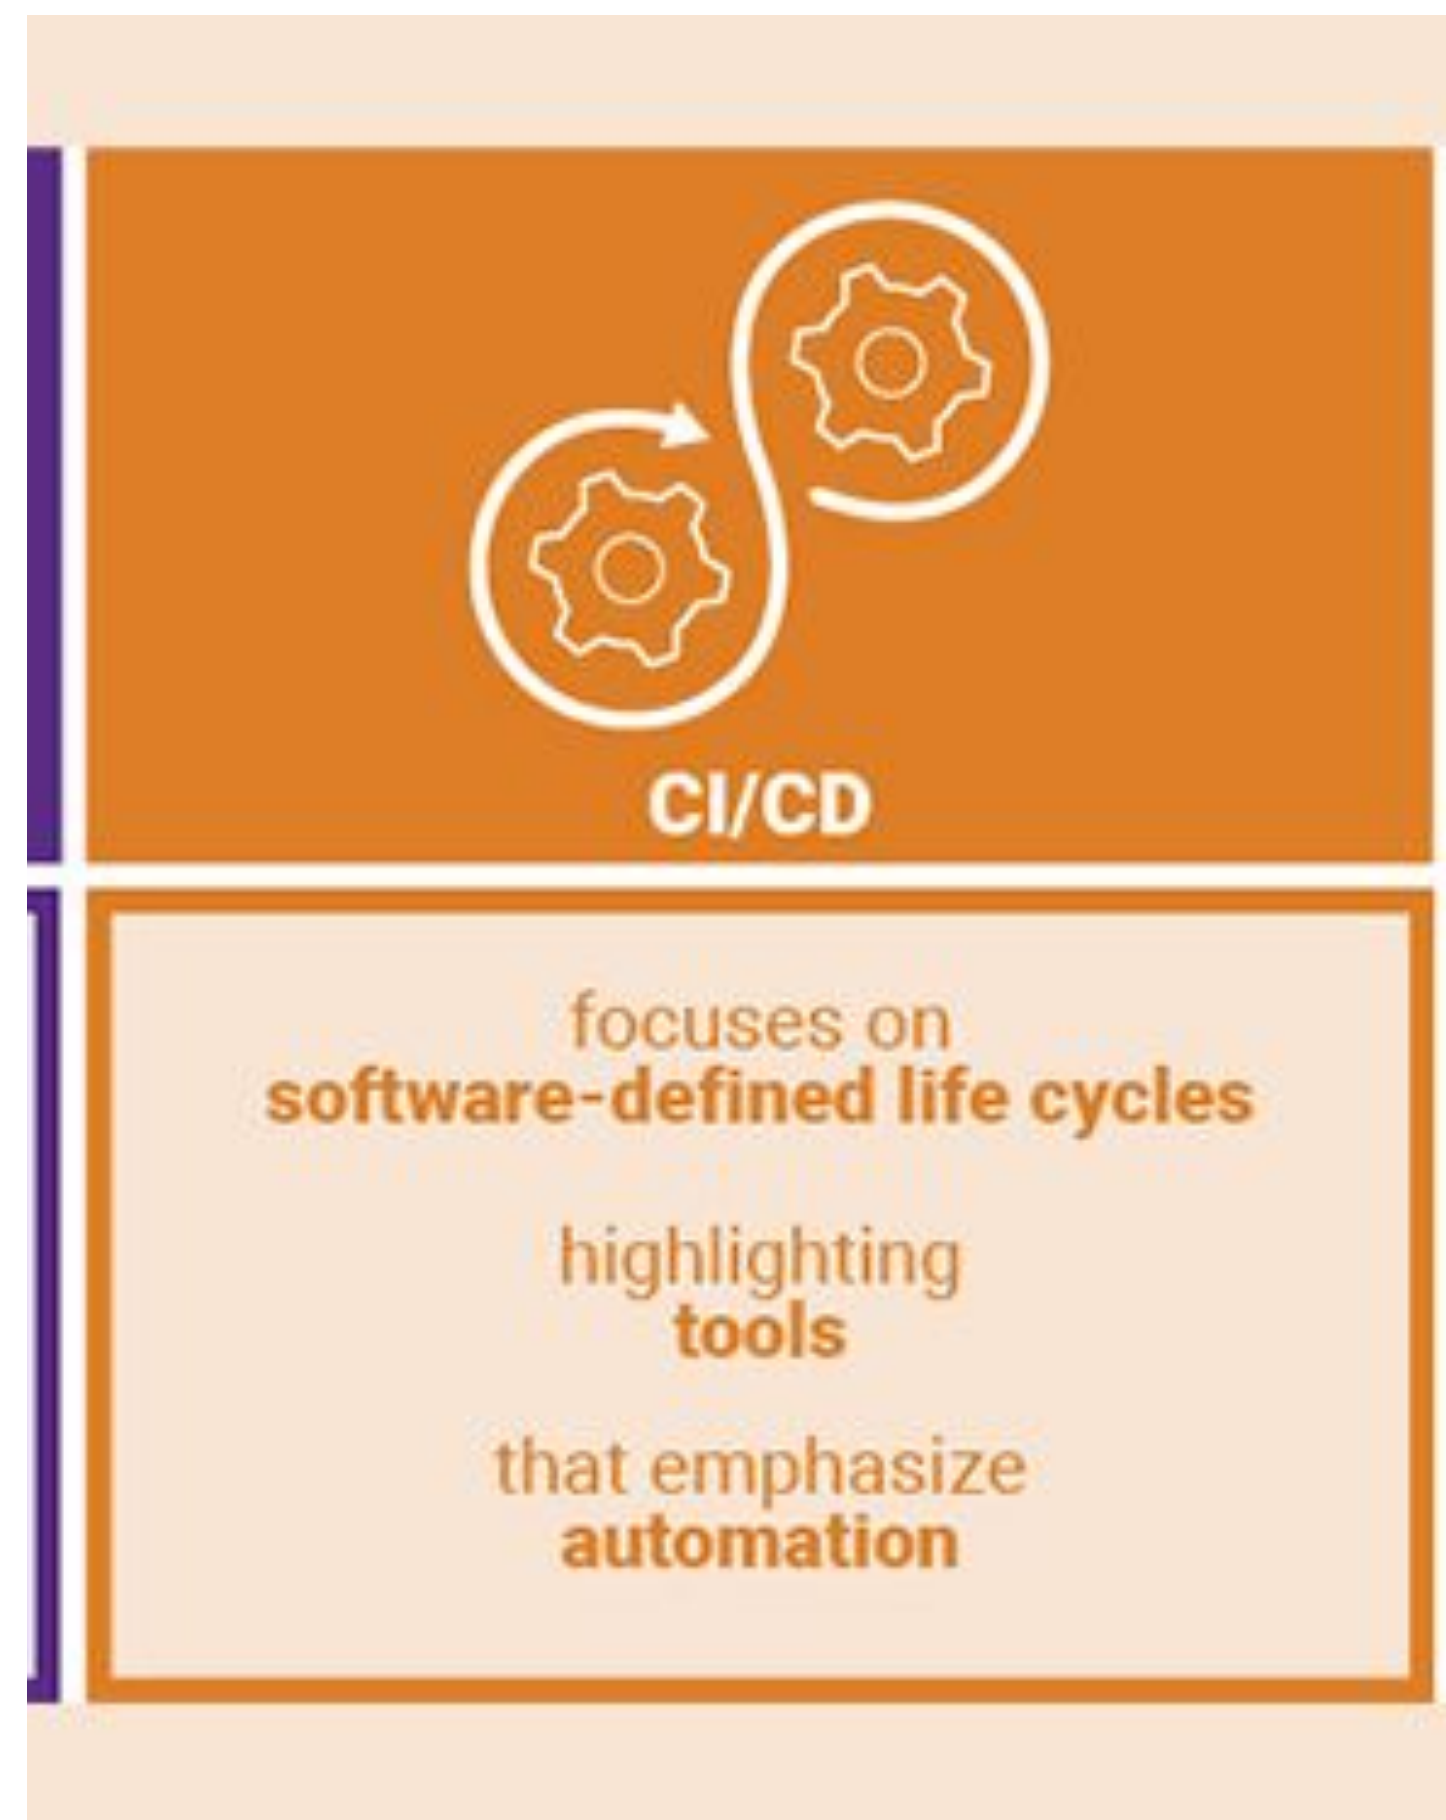
2019



Agility + CI/CD + DevOps = success




CI/CD Stakeholders



Developers




Operations
+
Security




QA
+
Product Officer
+
Users

CI/CD Stakeholder Concerns




Devs

- Don't write CI pipelines
- Don't like application manifests
- Don't like yaml!
- Want visibility into CI/CD process



Ops

- Don't like difficult to understand CI configuration
- Don't like difficult to reuse CI pipelines
- Don't like inflexible CI/CD
- Don't like yaml!
- Want K8s native CI pipelines
- Want robust application lifecycle management



QA/PO

- Don't like not knowing what version of which app they just tested out
- Don't like not having a list of all deployed applications
- Don't like incorrectly informing users which features are available on prod



2019



Issues BioBox had with existing tools

- Drone
 - Can achieve modular pipelines via jsonnet plugin...but jsonnet is unfamiliar to developers
 - There was alpha support for Kubernetes runtime, however not configurable from CI config...Drone internally was creating Jobs/Pods, *was later deprecated* - [drone/drone-runtime/issues/69](https://github.com/drone/drone-runtime/issues/69)
- GitLab
 - K8s GitLab runner a huge blackbox, don't want to maintain a fork, also different scope
 - Reusability via YAML DSL (".partial: &partial", "<<: *partial" !?) is annoying for Ops, difficult for Devs
- Tektoncd/pipeline
 - Was seen as alternative to Argo Workflows, which we already had operational experience with




2019



Objectives - *Things We Knew We Wanted*

- Developers should feel comfortable reading and writing CI configuration
- CI pipelines should be kept DRY and modularized
- Flexible CI/CD configuration
- Audit log from Git event to deployed resources
- CI/CD observability tooling consistent with primary application
- Manual, schedule, or event-based triggering of CI pipelines
- Support many 3rd party dependencies (many Dockerfiles)
- Preview application for all PRs, easily accessible to QA/PO



2019



Architecture overview: CI workflows

```
apiVersion: argoproj.io/v1alpha1
kind: Workflow
metadata:
  generateName: hello-world-parameters-
spec:
  # invoke the whalesay template with
  # "hello world" as the argument
  # to the message parameter
  entrypoint: whalesay
  arguments:
    parameters:
      - name: message
        value: hello world

  templates:
  - name: whalesay
    inputs:
      parameters:
        - name: message          # parameter declaration
    container:
      # run cowsay with that message input parameter as args
      image: docker/whalesay
      command: [cowsay]
      args: ["{{inputs.parameters.message}}"]
```

Argo Workflows

- Like Job on steroids
- Parameters, Artifacts (Git, S3, +)
- Linear sequence of steps
- DAG of steps
- Retry-able

Architecture overview: Git webhooks



GitLab webhooks

- Push
- Tag
- New branch
- MR
 - open/close/update
- MR comments

Argo Workflows

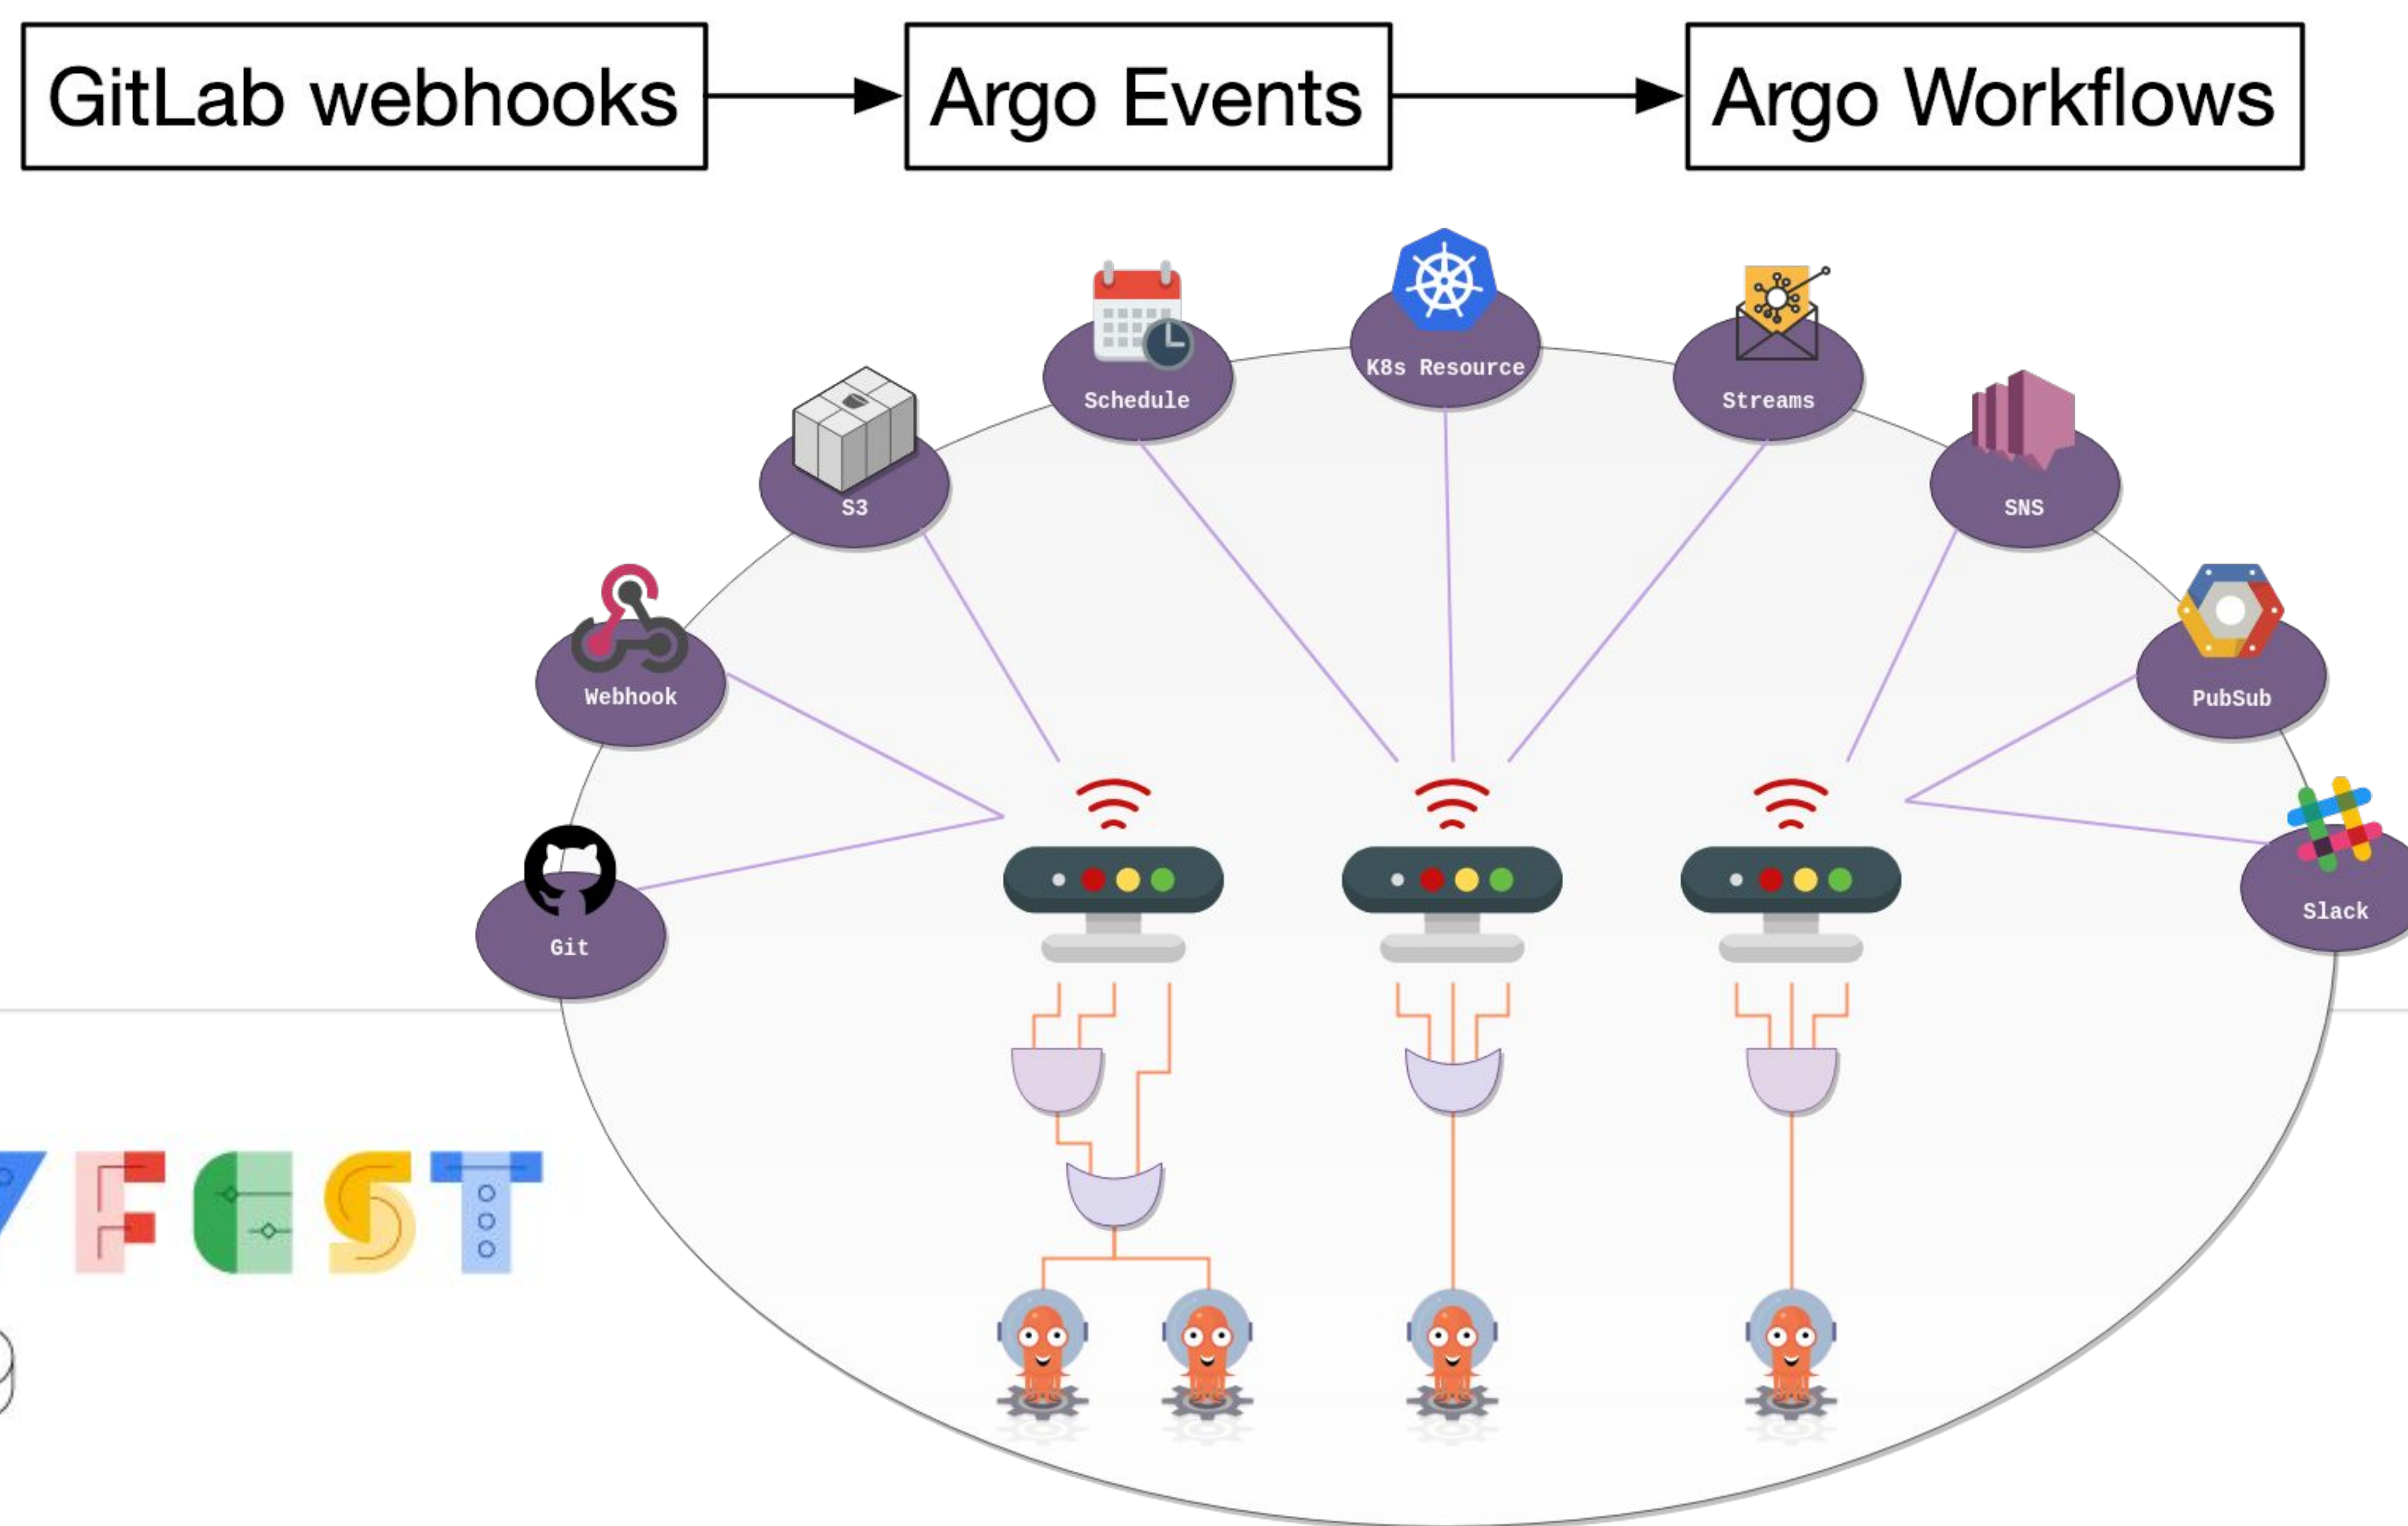


2019

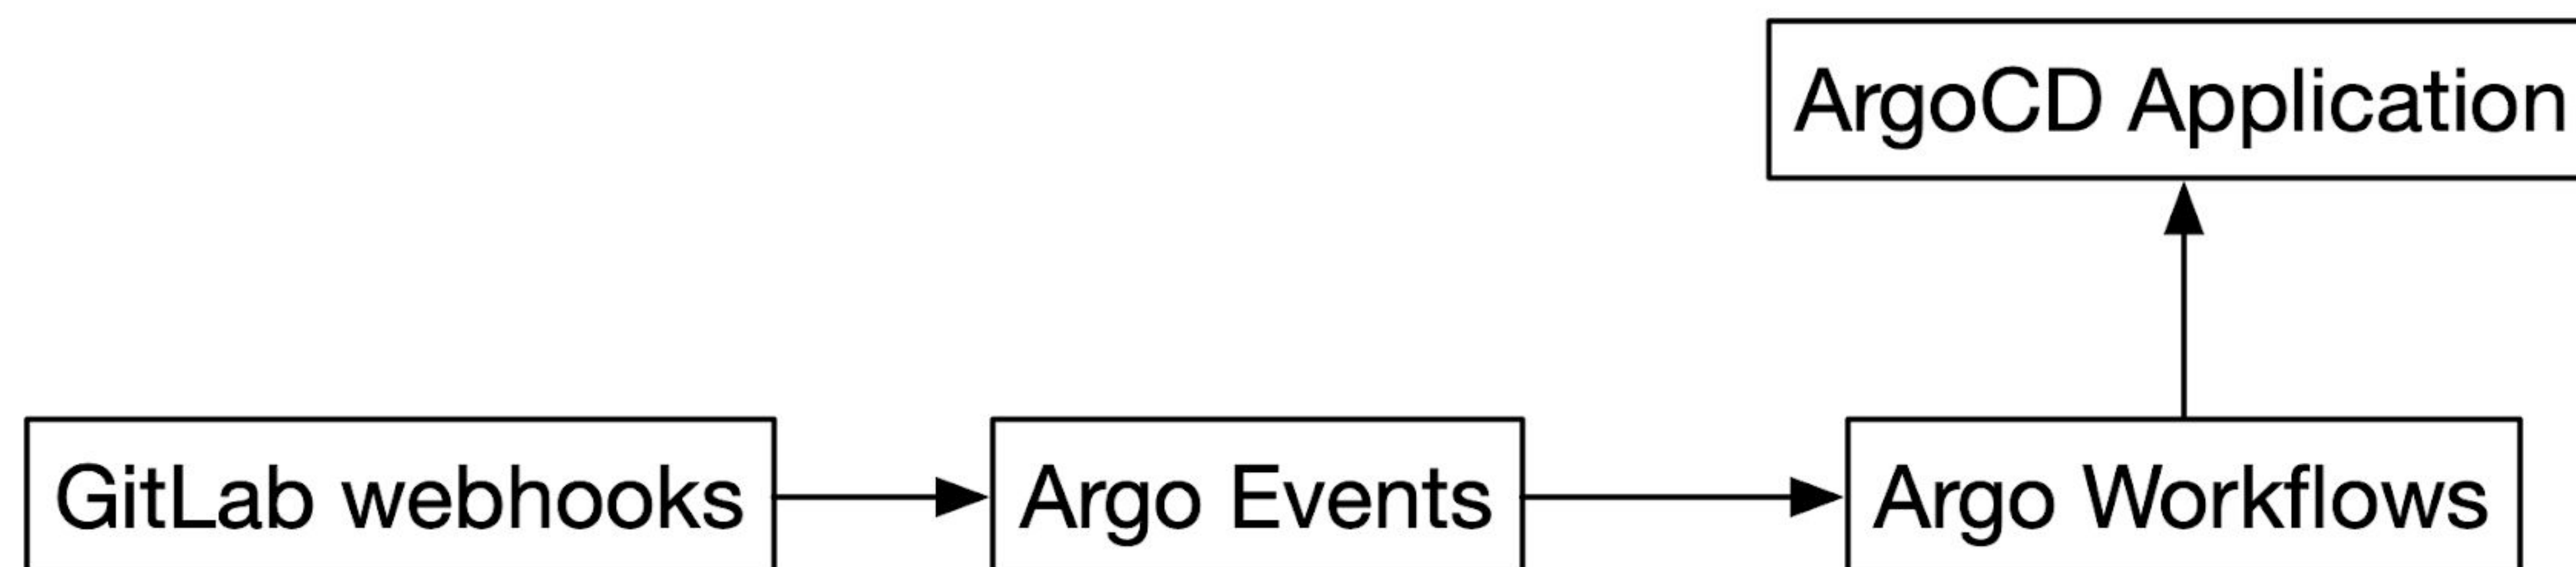


Architecture overview: Webhook to Workflow

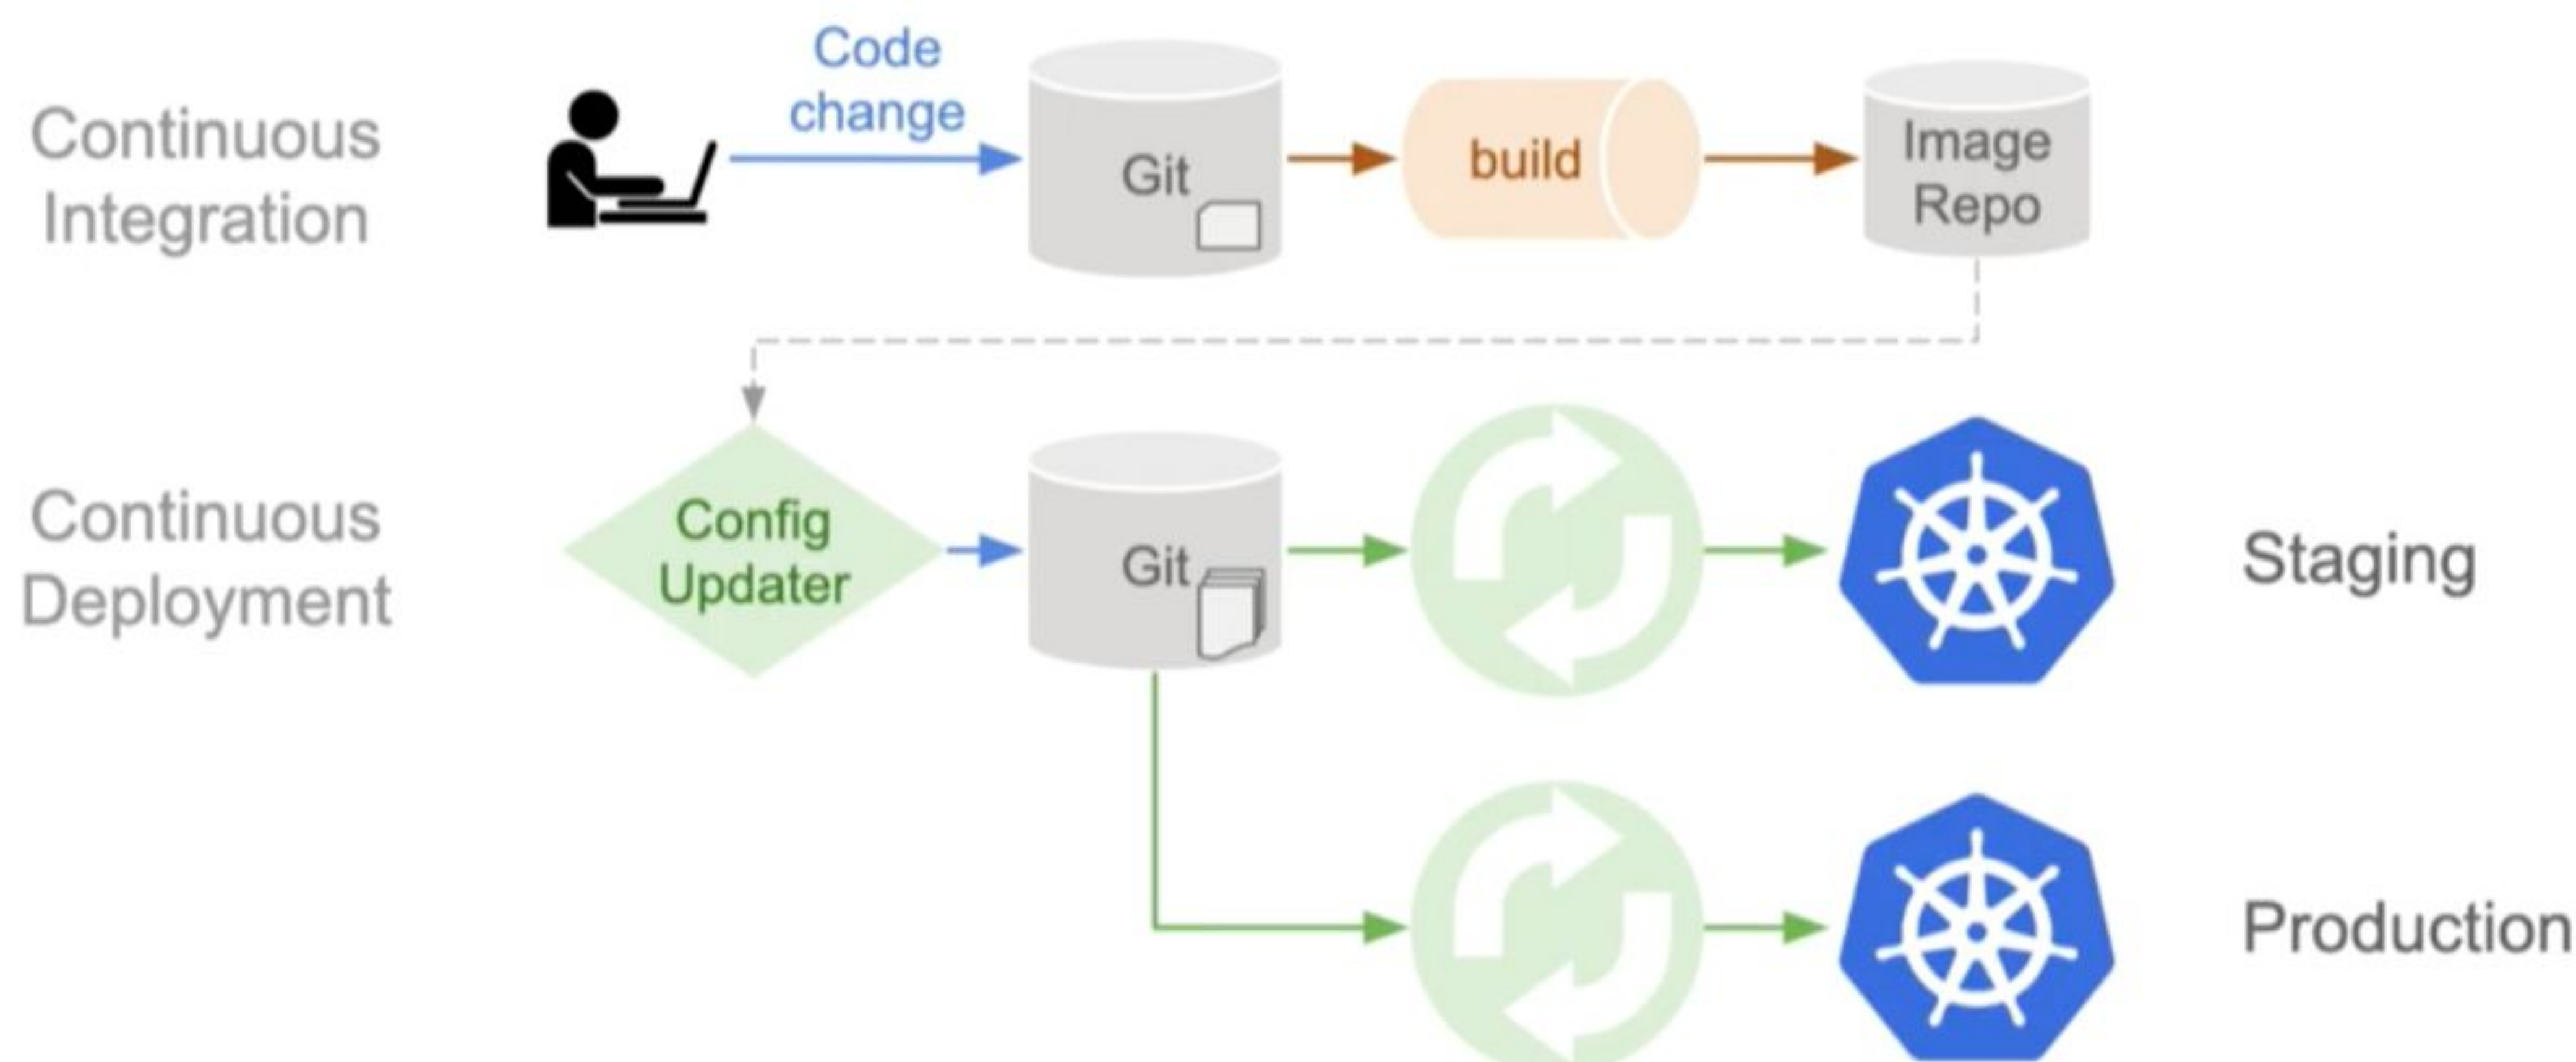
“Argo Events is an event-based dependency manager for Kubernetes which helps you define multiple dependencies from a variety of event sources like webhook, s3, schedules, streams etc. and trigger Kubernetes objects after successful event dependencies resolution.”



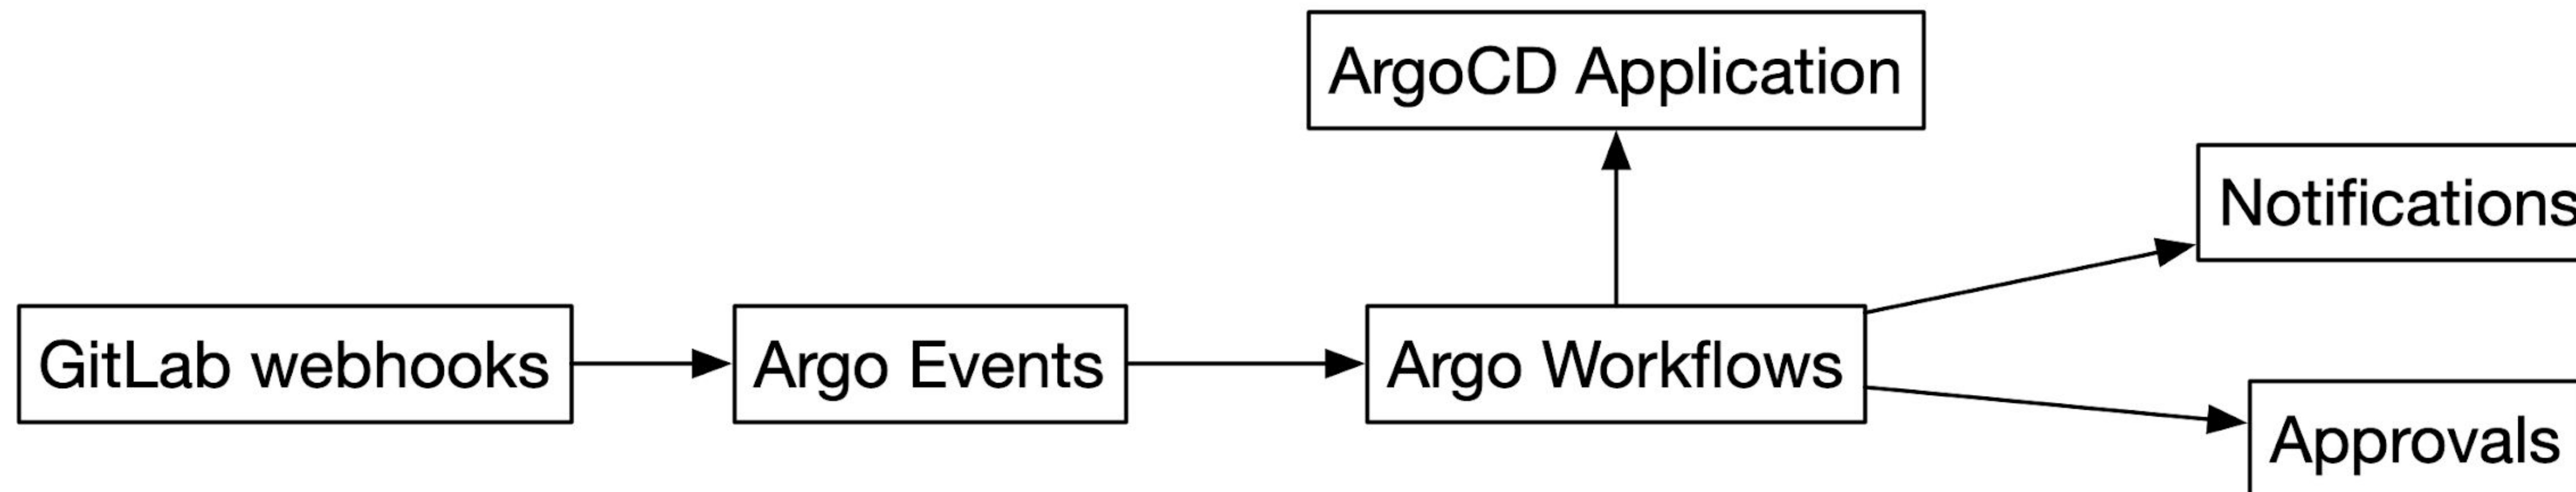
Architecture overview: GitOps “Application”



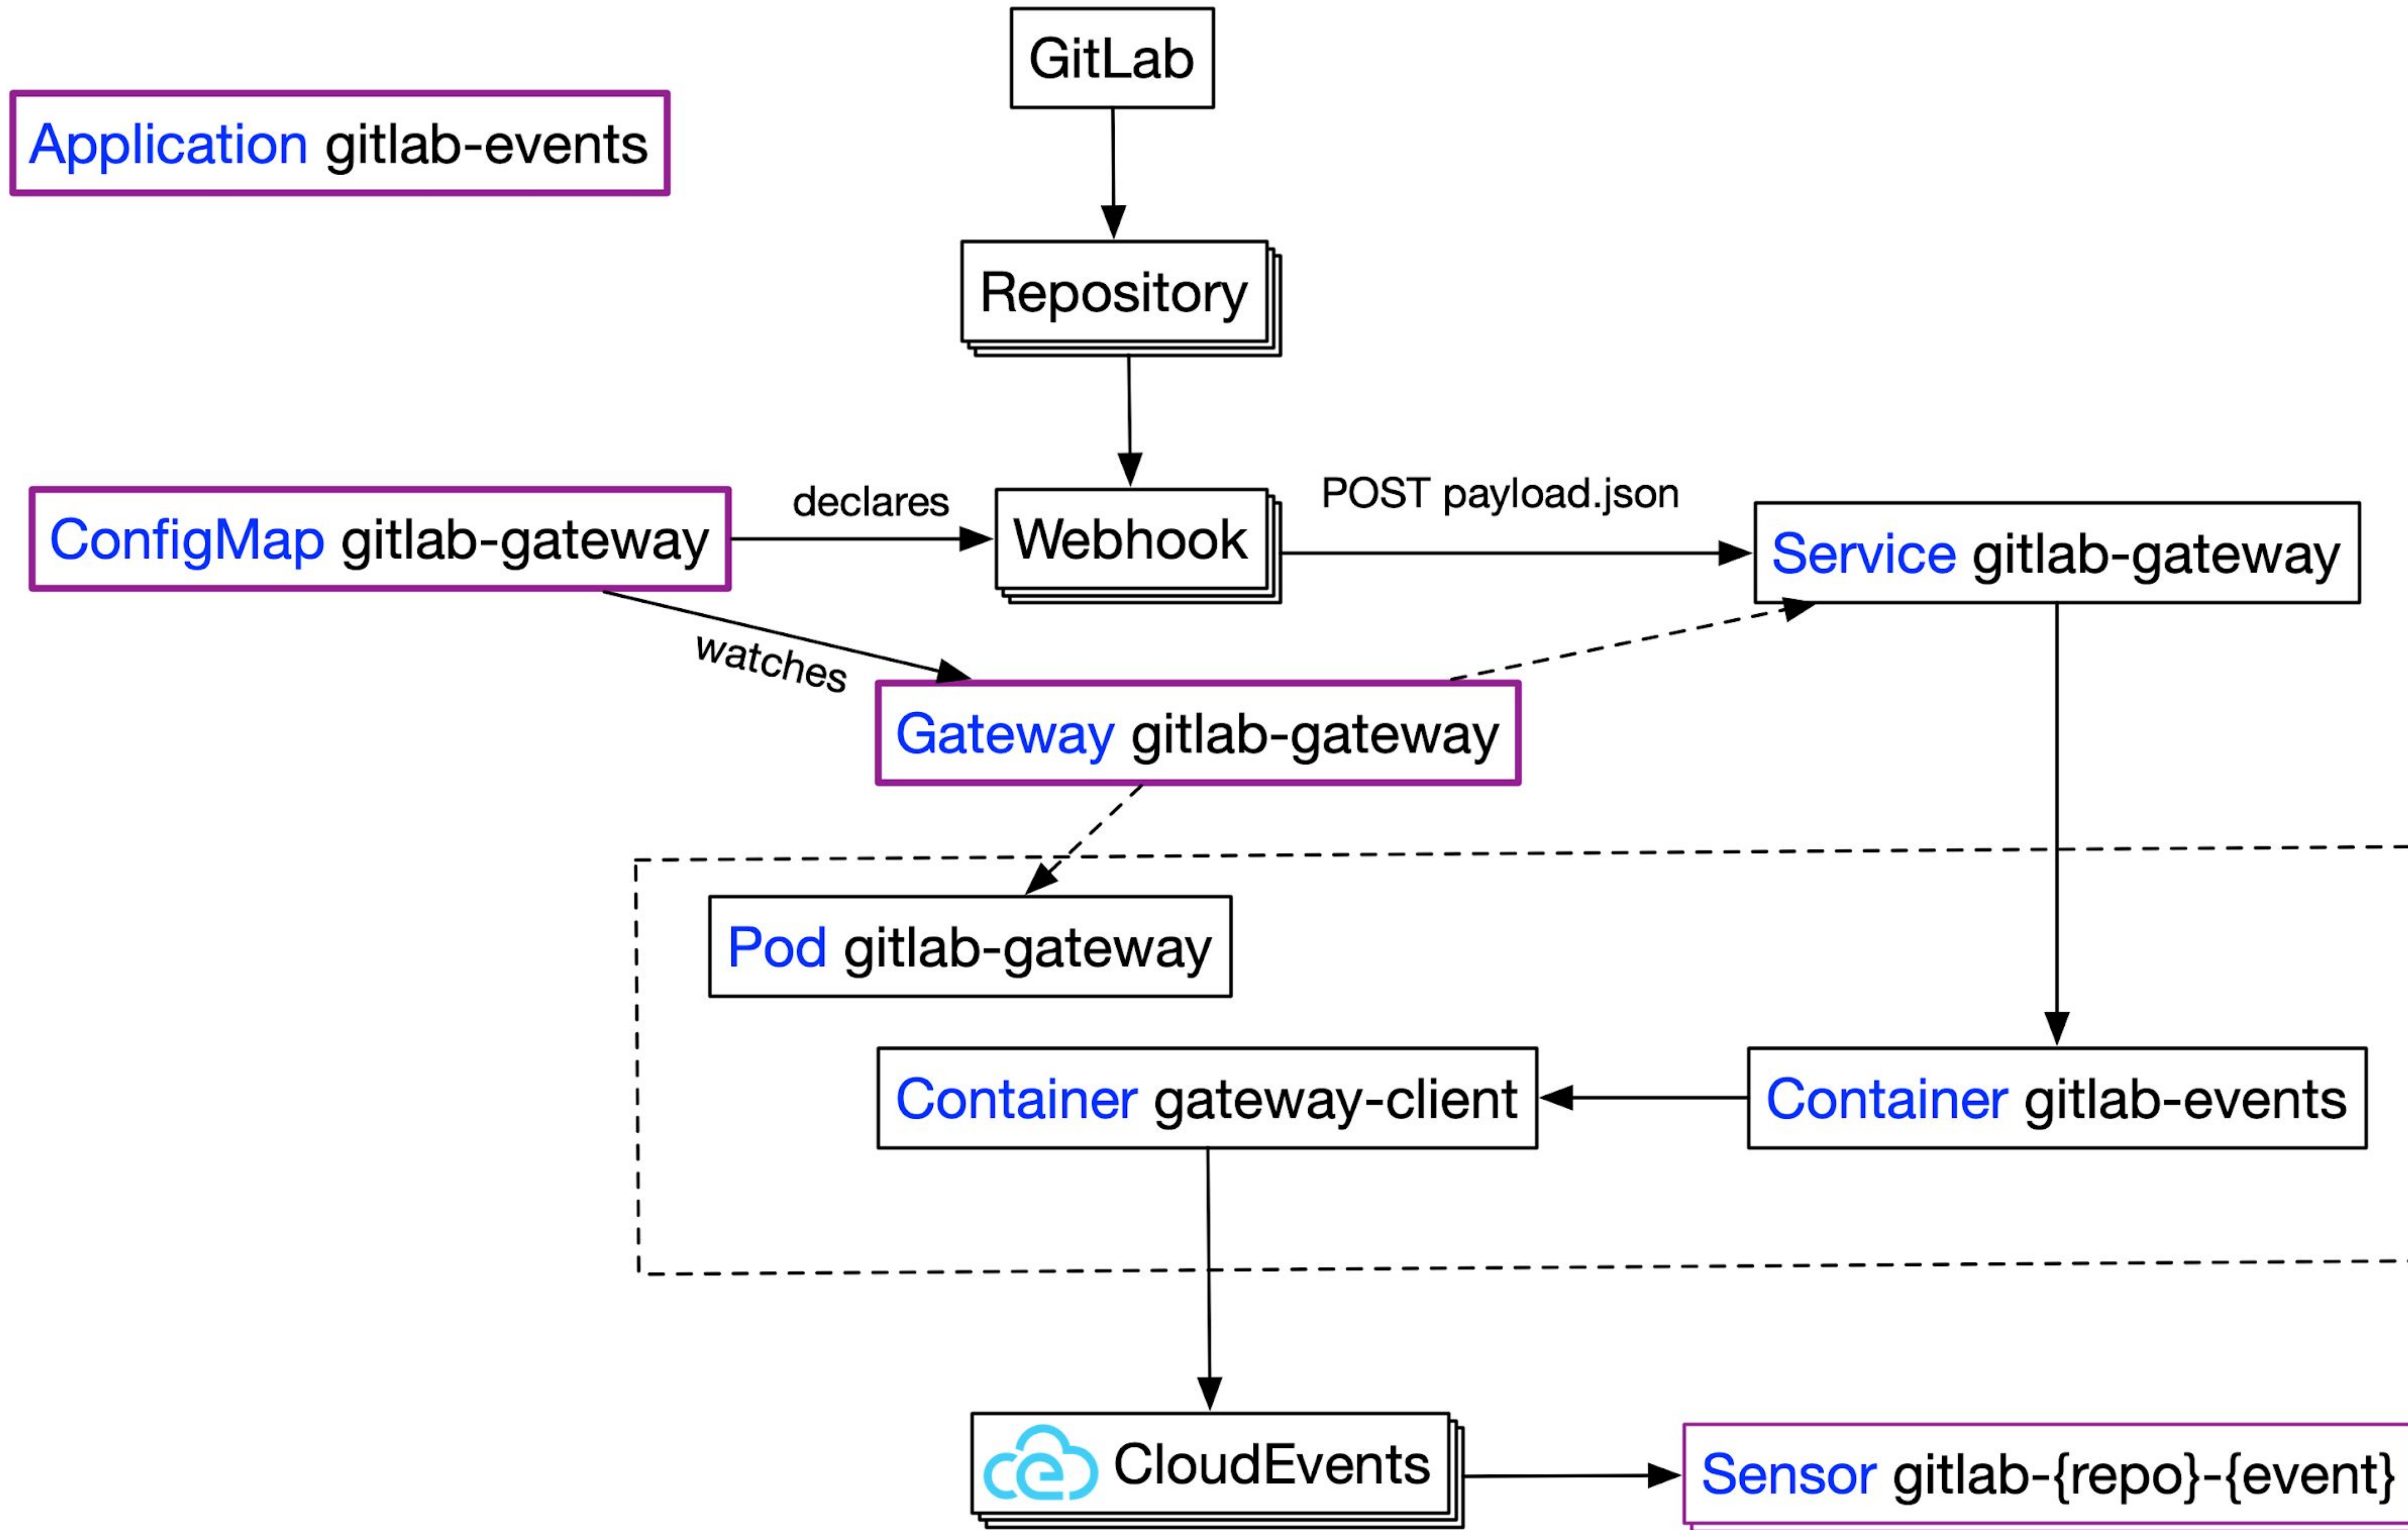
- GitOps “Application” CRD
- Defines git source and tracking
- Defines destination cluster
- Optionally defines tool settings - helm values, jsonnet top-level params, etc



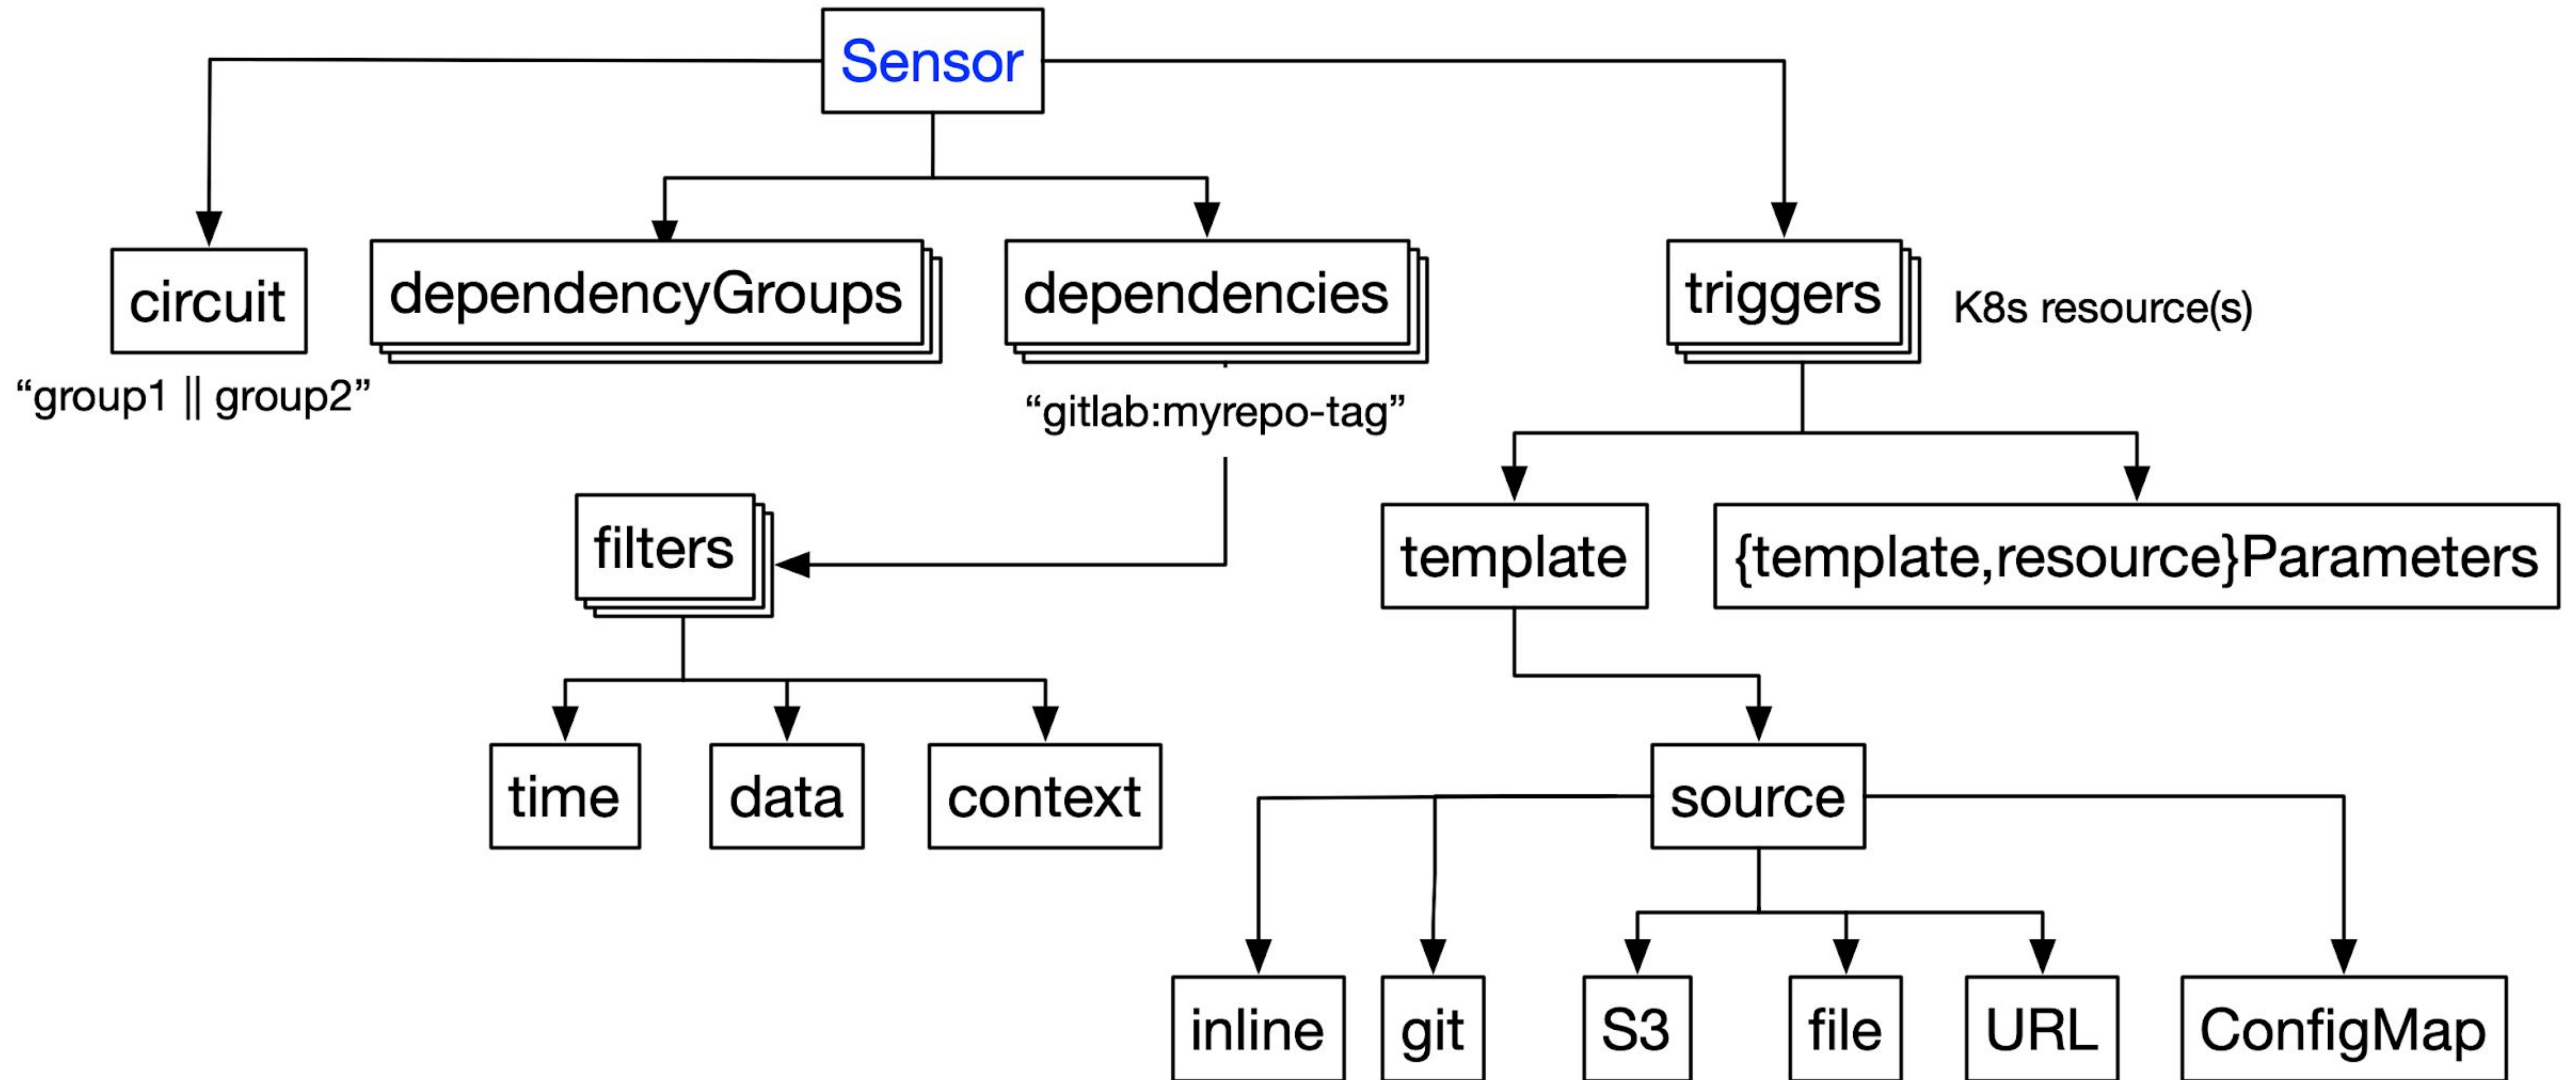
Architecture overview: What **is** a CI/CD software?




Technical Architecture: Gateway and Sensors



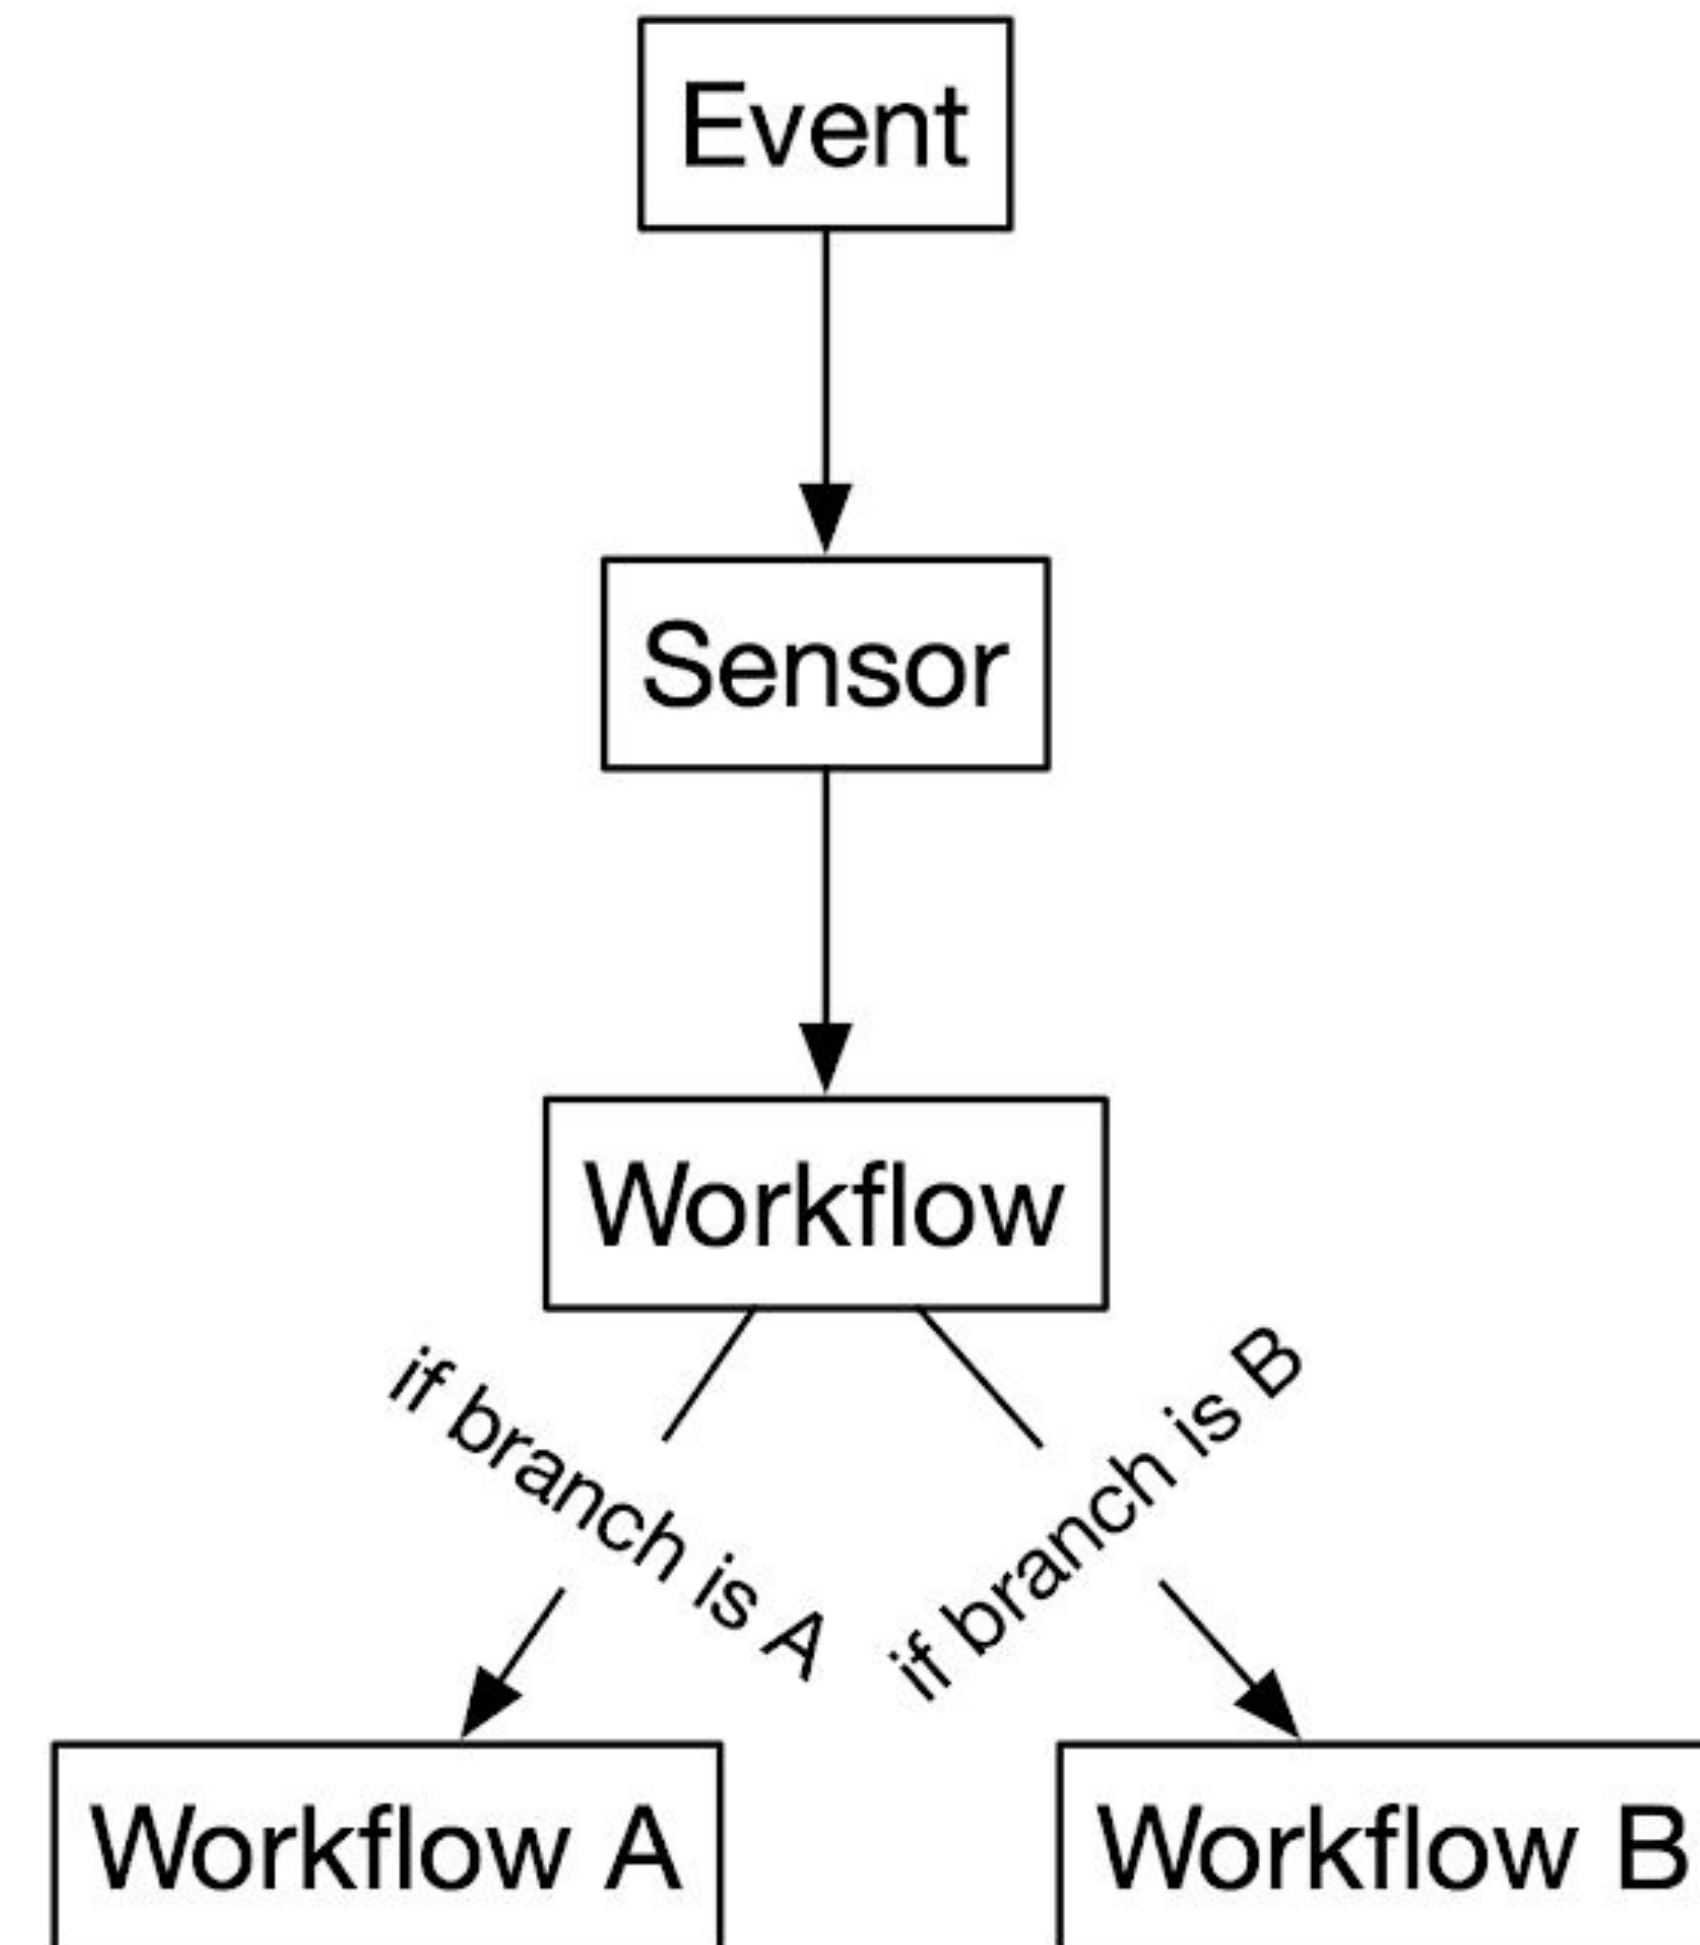
Sensor Spec



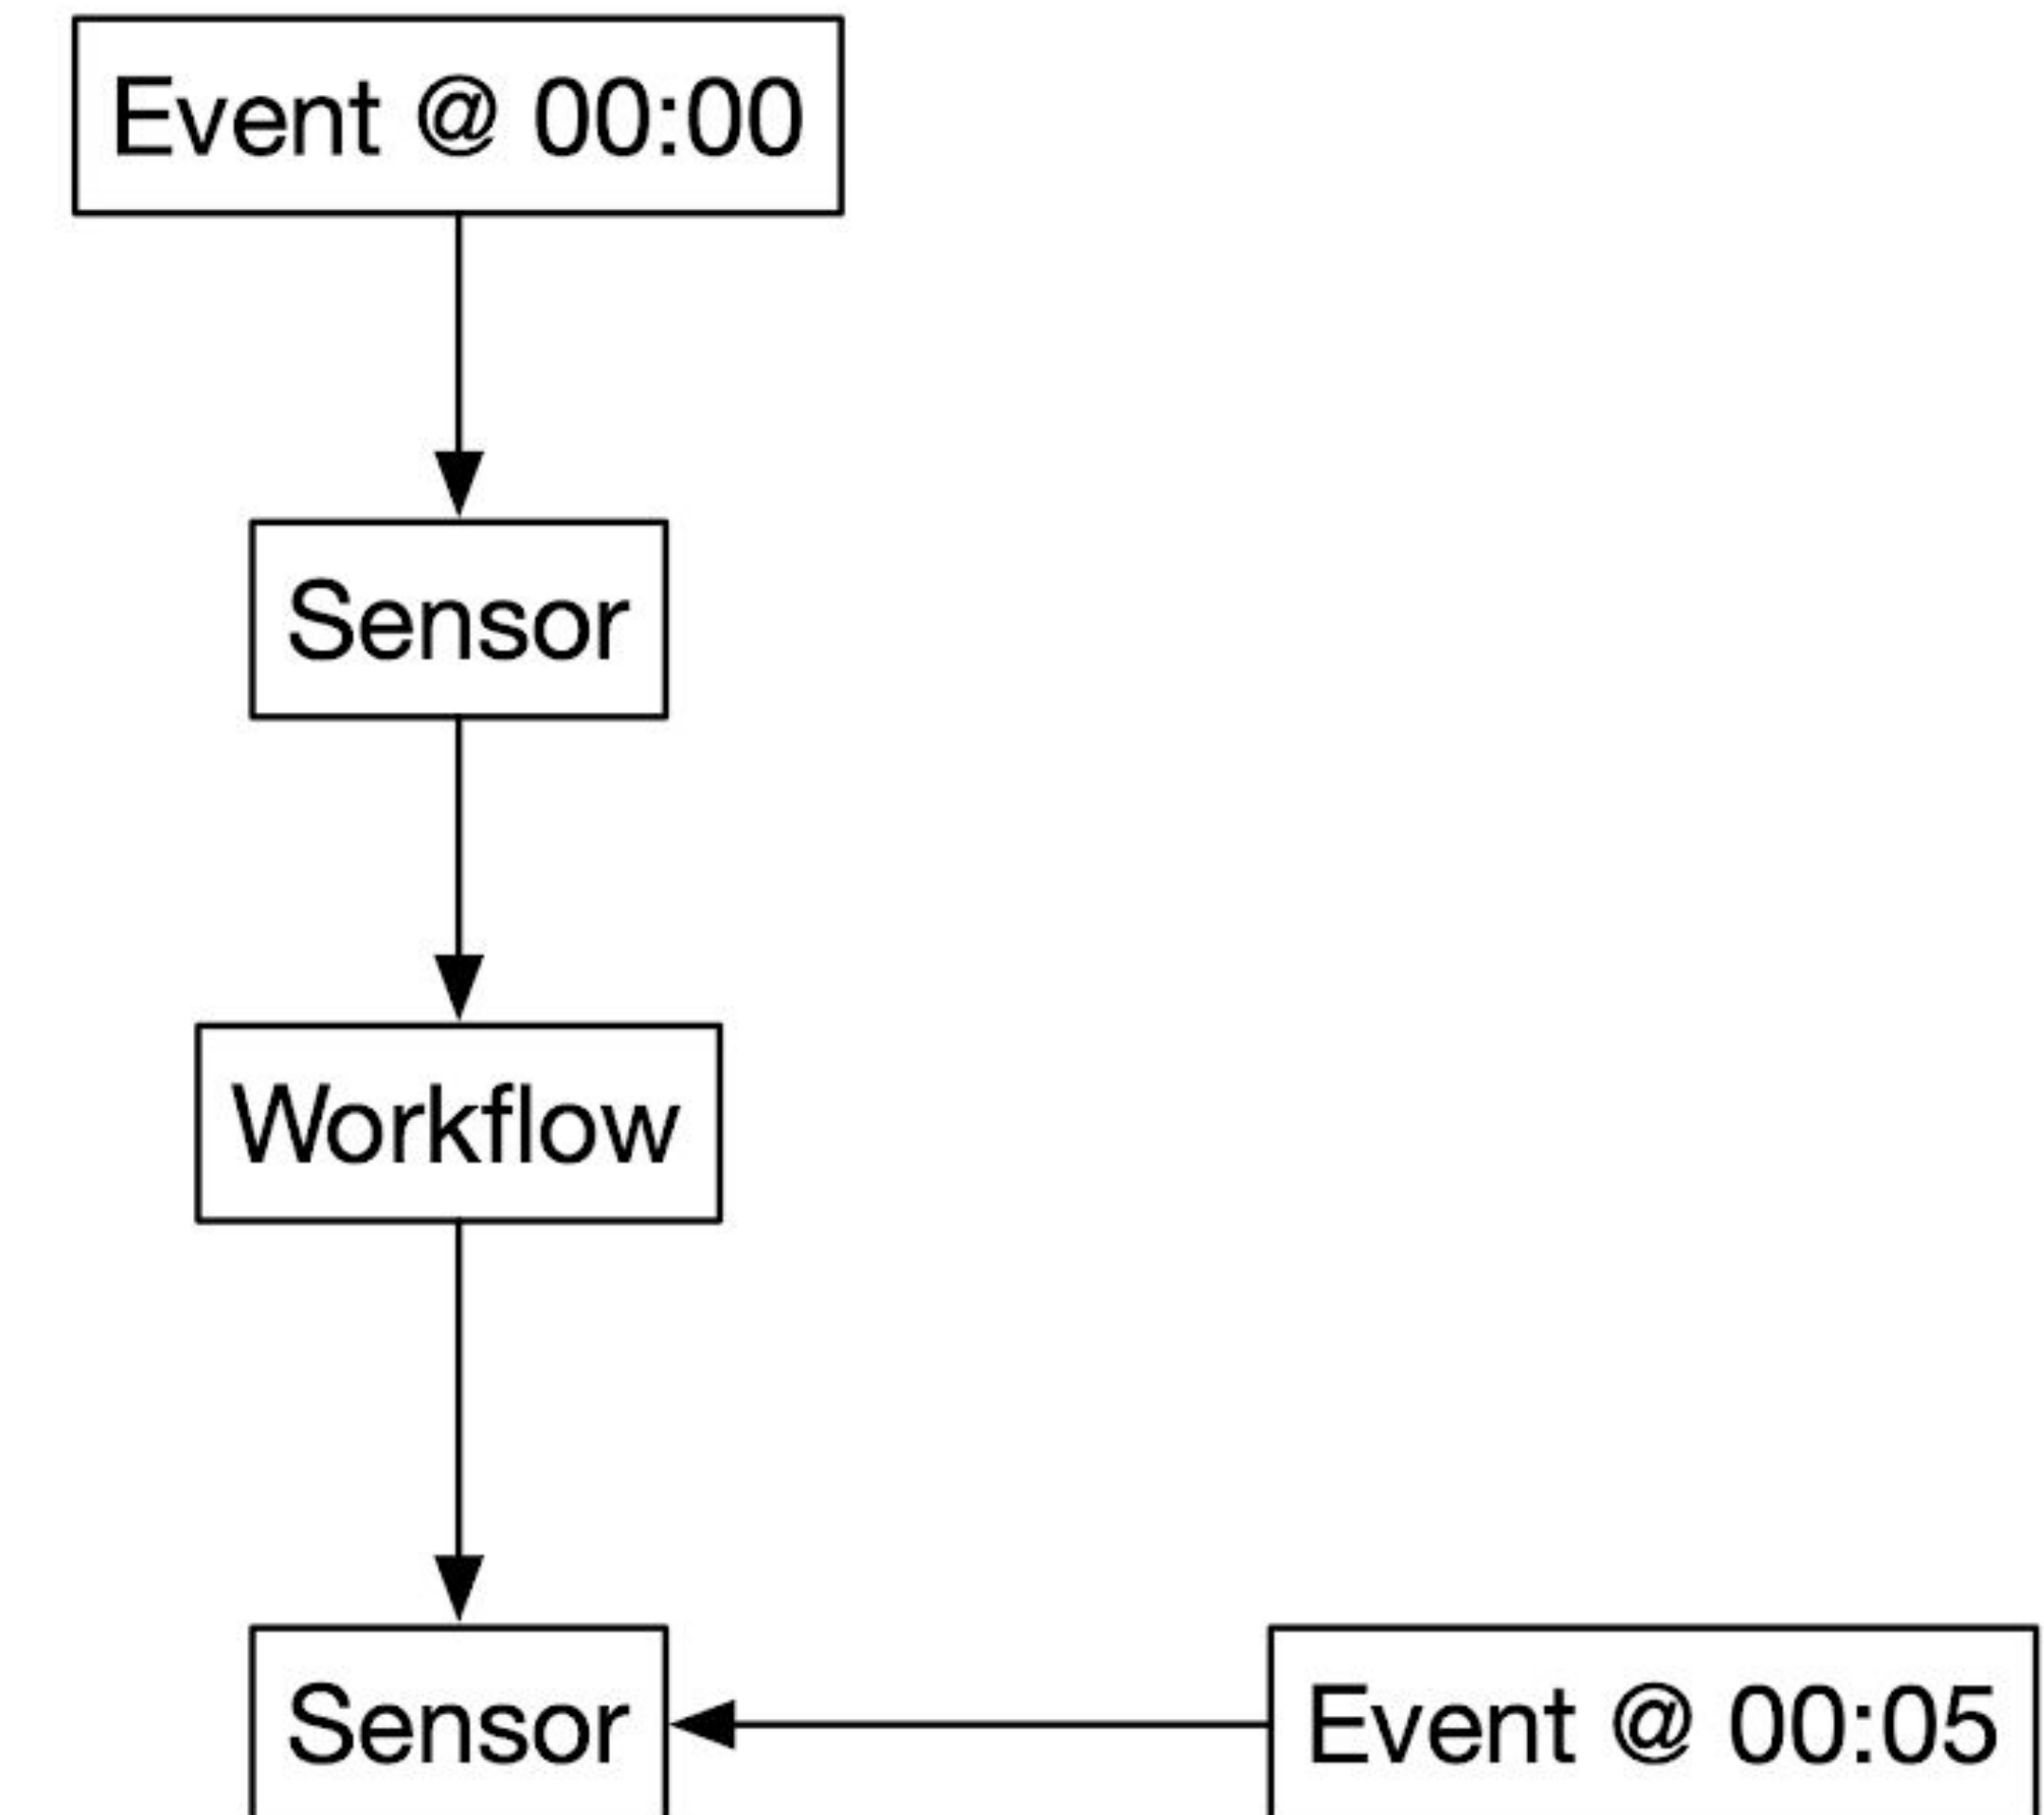
Sensor Designs



*logic contained within single Workflow,
CI pipeline runs in single Workflow*



*logic contained within single Workflow,
parent Workflow programmatically generates children Workflows*



logic contained within filters of Sensor



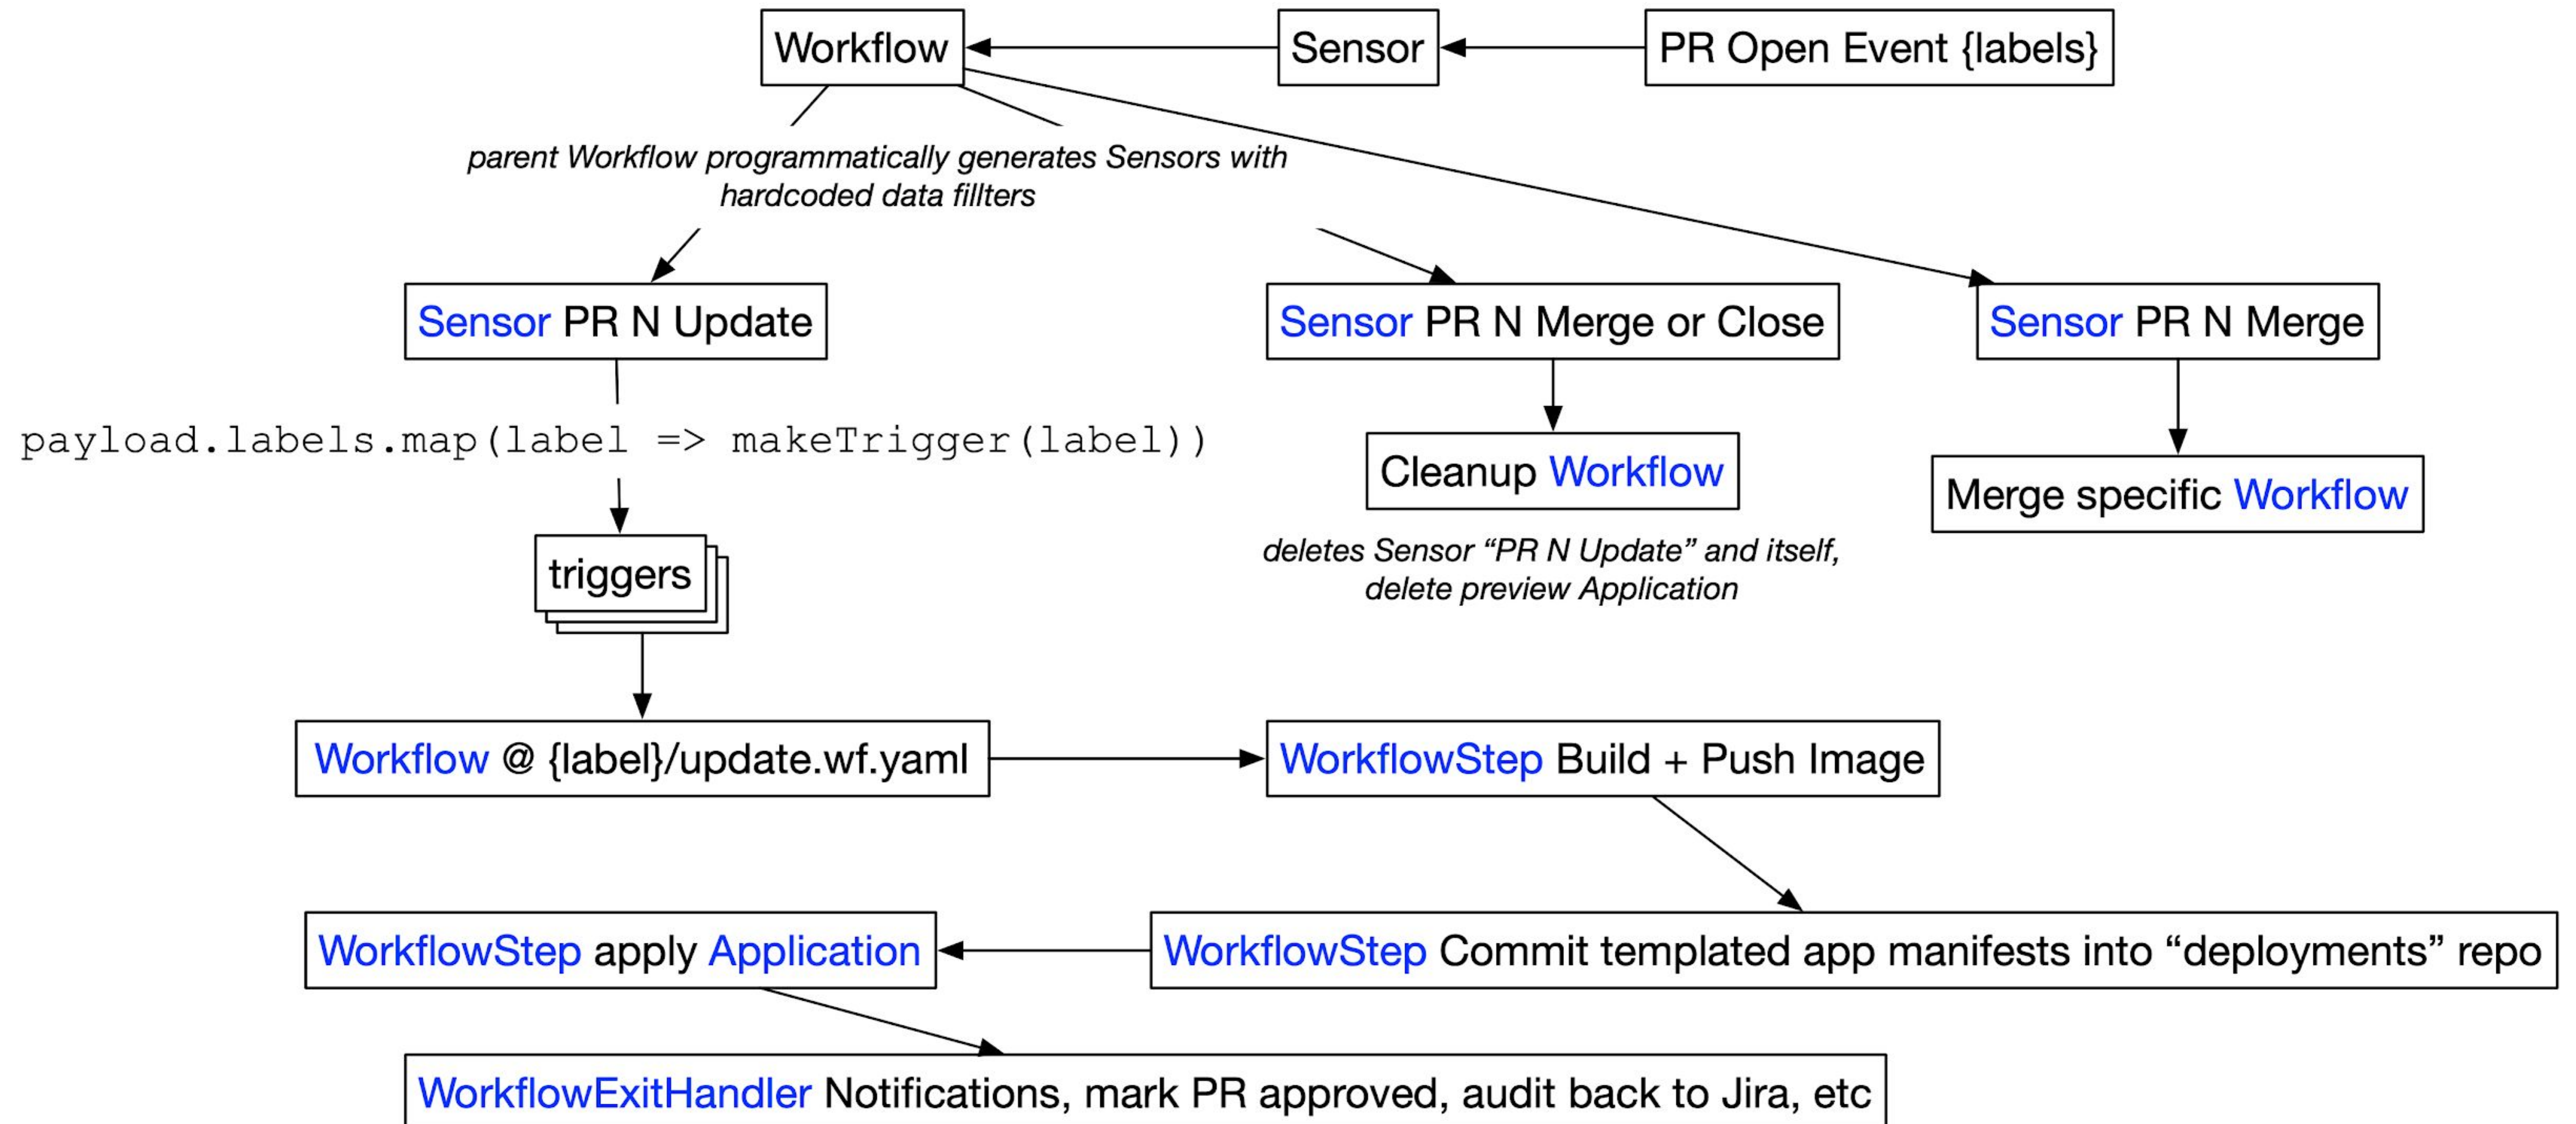
2019

```
spec:
  dependencies:
    - filters:
      data:
        - path: object_attributes.iid
          type: number
          value:
            - "31"
        - path: object_attributes.action
          type: string
          value:
            - update
      name: merge-request-update
      name: gitlab-gateway:bx-merge-request
```



BioBox Monorepo CI/CD

- Each merge request is annotated with labels specifying which services to deploy in a monorepo
- Developer can test one or more altered services in the context of the entire stack (the rest deployed from whichever was latest release)



Results

- CI logic can be written in any language that developers are comfortable with, breaking down divisional roles between Dev and Ops
- Kubernetes CI Workflows can be labeled/annotated with repo/branch/tag etc (templateParameters)
- GitOps for CD via ArgoCD enables visibility for QA/PO as well as robustness for Ops
- Flexibility - receive a webhook (or event!), code process payload, pick Sensor design that fits task
- Reuse of observability stack for metrics and logging on CI workflows
- CI workflows autoscaling via K8s resource requests, scheduling via tolerations and node taints
- Consistent tooling - developers can get familiar with K8s through CI/CD, same K8s for primary app
- Arbitrary notifications (slack, PR comments, email) written as Argo Workflow steps
- CI workflows can be manually triggered via kubectl/argo CLI, or by Argo Events (e.g. GCR PubSub)



2019



Future Objectives

- Improve multi-event multi-sensor Workflow visibility
- Argo Workflows still in YAML... working on K8s TypeScript client implementation
- Argo Workflows 2.4 release will bring “Template CRD” - reuse workflow steps across Workflows!
- Argo Events has support for NATS instead of HTTP streaming for Gateway
 - Kafka is a supported event source, but have to use NATS for Gateway-Sensor streaming?
 - Event replayability, long term storage, audit logging
- Special pipelines for PRs which are WIP - bring up web IDE, run apps in debug mode




2019



Questions & demo!

Thank you!

- Intuit + Blackrock
- Argo slack
- Devfest organizers!



2019

