



KubeCon



CloudNativeCon

Europe 2022

WELCOME TO VALENCIA





KubeCon



CloudNativeCon

Europe 2022

TikTok's Story:

How To Manage a Thousand Applications on Edge With Argo CD

Qingkun Li - TikTok/Bytedance, Inc.

Jesse Suen - Akuity, Inc.



TikTok's Story:

How To Manage a Thousand Applications on Edge With Argo CD

This talk describes a case study of how TikTok manages 3000 applications across 100 global edge clusters with Argo CD. We discuss considerations, tips, and techniques for using Argo CD to manage cluster applications on the edge.



Qingkun Li
Tech Lead Manager
TikTok/Bytedance, Inc.



Jesse Suen
Co-Founder & CTO
Akuity, Inc.



TikTok's Edge Cluster Overview



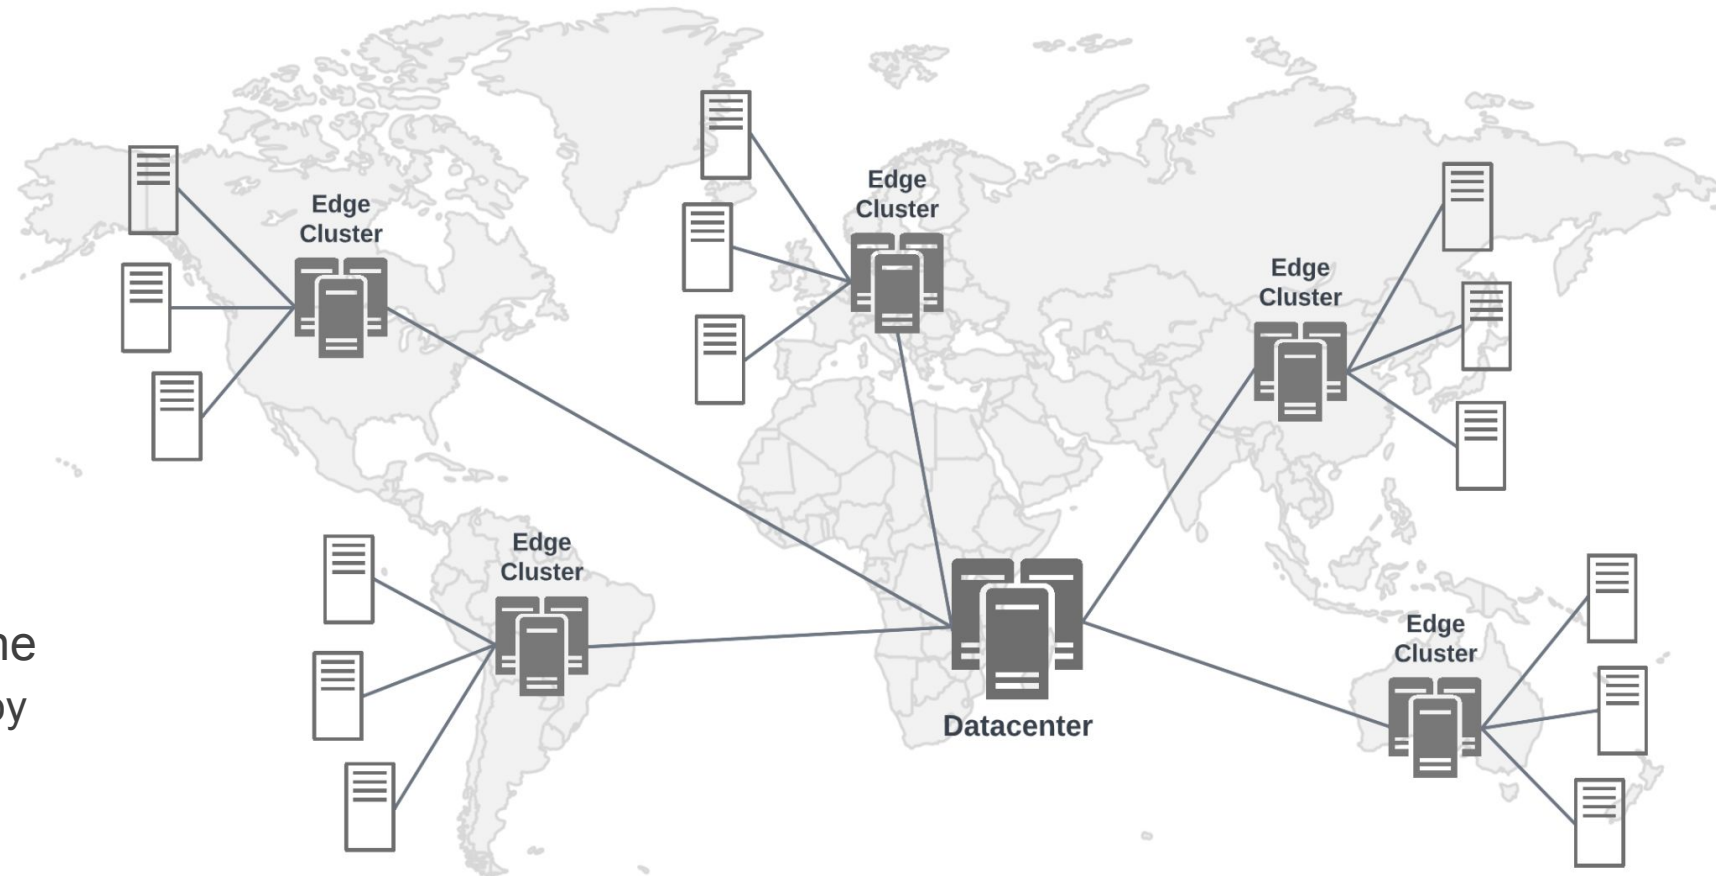
KubeCon



CloudNativeCon

Europe 2022

- Distributed around the world
 - ~100 edge clusters
 - 10-60 nodes on each
- TikTok edge services
 - video CDN cache
 - live streaming
 - game
 - etc.
- Datacenter
 - Management control plane
 - e.g. edge service deploy
 - Service data plane
 - e.g. CDN origin



TikTok's Edge Cluster Deployment



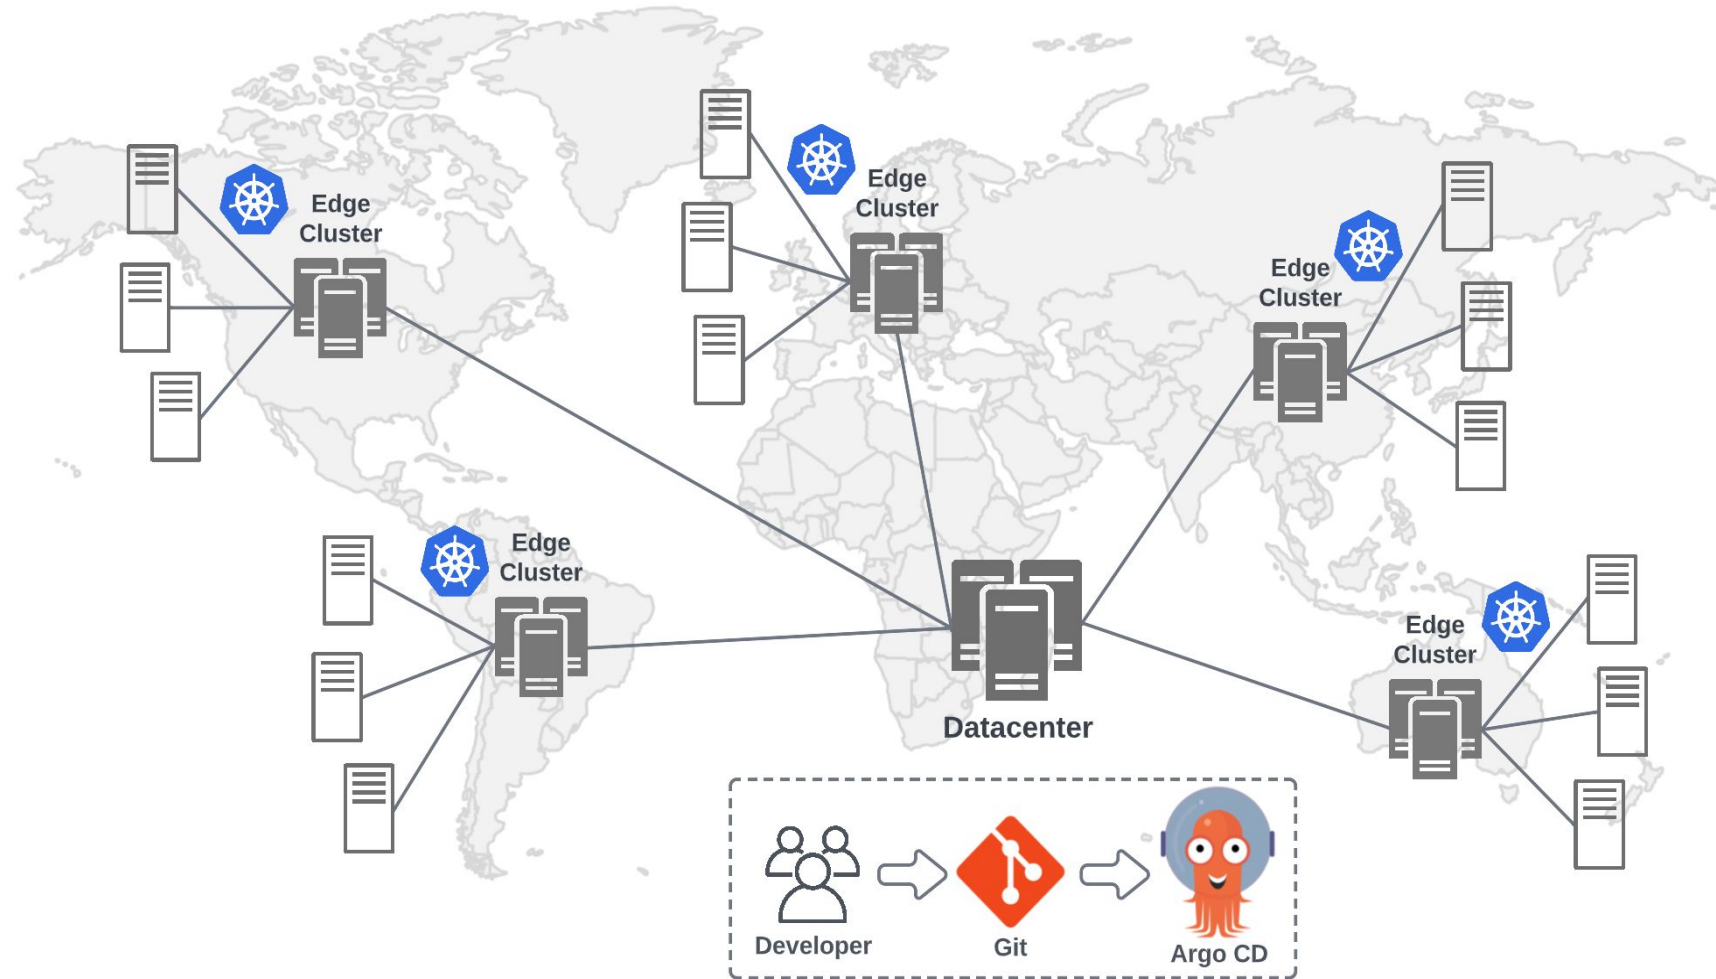
KubeCon



CloudNativeCon

Europe 2022

- Each edge is a K8s cluster
- GitOps from datacenter
 1. Developer push K8s configuration to git
 2. Central Argo CD controller pulls config from git and sync to edge K8s clusters



Edge Deployment Pattern

- An edge service is deployed to many edge clusters
- The service functionality and behavior are similar on all edge clusters
- As a result, the K8s configuration pattern of an edge service:
 - Deployment on all edge clusters shares large common configurations
 - There are some cluster-specific configurations
 - replica count
 - resource quota
 - IP address / ingress host
 - etc.

Edge Deployment Pattern - Example



KubeCon



CloudNativeCon

Europe 2022

Edge Cluster 1

```
1 apiVersion: apps/v1
2 kind: Deployment
3 metadata:
4   name: nginx
5 spec:
6   replicas: 3
7   selector:
8     matchLabels:
9       app: nginx
10  template:
11    metadata:
12      labels:
13        app: nginx
14    spec:
15      containers:
16      - image: nginx:1.19.0
17        name: nginx
18        resources:
19          requests:
20            cpu: "100m"
21          limits:
22            memory: "256Mi"
23            cpu: "150m"
24
25 ---
26
27 apiVersion: v1
28 kind: Service
29 metadata:
30   name: nginx
31 spec:
32   selector:
33     app: nginx
34   externalIPs:
35   - 1.1.1.1
36   ports:
37   - name: http
38     protocol: TCP
39     port: 80
```

Edge Cluster 2

```
1 apiVersion: apps/v1
2 kind: Deployment
3 metadata:
4   name: nginx
5 spec:
6   replicas: 6
7   selector:
8     matchLabels:
9       app: nginx
10  template:
11    metadata:
12      labels:
13        app: nginx
14    spec:
15      containers:
16      - image: nginx:1.19.0
17        name: nginx
18        resources:
19          requests:
20            cpu: "100m"
21          limits:
22            memory: "256Mi"
23            cpu: "150m"
24
25 ---
26
27 apiVersion: v1
28 kind: Service
29 metadata:
30   name: nginx
31 spec:
32   selector:
33     app: nginx
34   externalIPs:
35   - 2.2.2.2
36   ports:
37   - name: http
38     protocol: TCP
39     port: 80
```

Edge Cluster 3

```
1 apiVersion: apps/v1
2 kind: Deployment
3 metadata:
4   name: nginx
5 spec:
6   replicas: 2
7   selector:
8     matchLabels:
9       app: nginx
10  template:
11    metadata:
12      labels:
13        app: nginx
14    spec:
15      containers:
16      - image: nginx:1.20.0
17        name: nginx
18        resources:
19          requests:
20            cpu: "100m"
21          limits:
22            memory: "256Mi"
23            cpu: "150m"
24
25 ---
26
27 apiVersion: v1
28 kind: Service
29 metadata:
30   name: nginx
31 spec:
32   selector:
33     app: nginx
34   externalIPs:
35   - 3.3.3.3
36   ports:
37   - name: http
38     protocol: TCP
39     port: 80
```

Edge Deployment Management - Kustomize



KubeCon



CloudNativeCon

Europe 2022

- Base Layer Configuration

```
deploy/  
├─ base/  
│   ├── kustomization.yaml  
│   ├── deployment.yaml  
│   └── service.yaml  
└─ overlays/  
    ├── cluster1/  
    │   ├── kustomization.yaml  
    │   ├── deployment.yaml  
    │   └── service.yaml  
    ├── cluster2/  
    │   ├── kustomization.yaml  
    │   ├── deployment.yaml  
    │   └── service.yaml  
    └── cluster3/  
        ├── kustomization.yaml  
        ├── deployment.yaml  
        └── service.yaml
```

```
1 resources:  
2 - deployment.yaml  
3 - service.yaml
```

```
1 apiVersion: apps/v1  
2 kind: Deployment  
3 metadata:  
4   name: nginx  
5 spec:  
6   replicas: 3  
7   selector:  
8     matchLabels:  
9       app: nginx  
10  template:  
11    metadata:  
12      labels:  
13        app: nginx  
14  spec:  
15    containers:  
16      - image: nginx:1.19.0  
17        name: nginx  
18        resources:  
19          requests:  
20            cpu: "100m"  
21          limits:  
22            memory: "256Mi"  
23            cpu: "150m"
```

```
1 apiVersion: v1  
2 kind: Service  
3 metadata:  
4   name: nginx  
5 spec:  
6   selector:  
7     app: nginx  
8   ports:  
9     - name: http  
10       protocol: TCP  
11       port: 80
```


Edge Deployment Management - Kustomize



KubeCon



CloudNativeCon

Europe 2022

- Overlay Configuration

deploy/

|— base/

| |— kustomization.yaml

| |— deployment.yaml

| |— service.yaml

|— overlays/

| |— cluster1/

| | |— kustomization.yaml

| | |— deployment.yaml

| | |— service.yaml

| |— cluster2/

| | |— kustomization.yaml

| | |— deployment.yaml

| | |— service.yaml

| |— cluster3/

| | |— kustomization.yaml

| | |— deployment.yaml

| | |— service.yaml

```
1 resources:
2 - ../../base
3 patchesStrategicMerge:
4 - deployment.yaml
5 - service.yaml
```

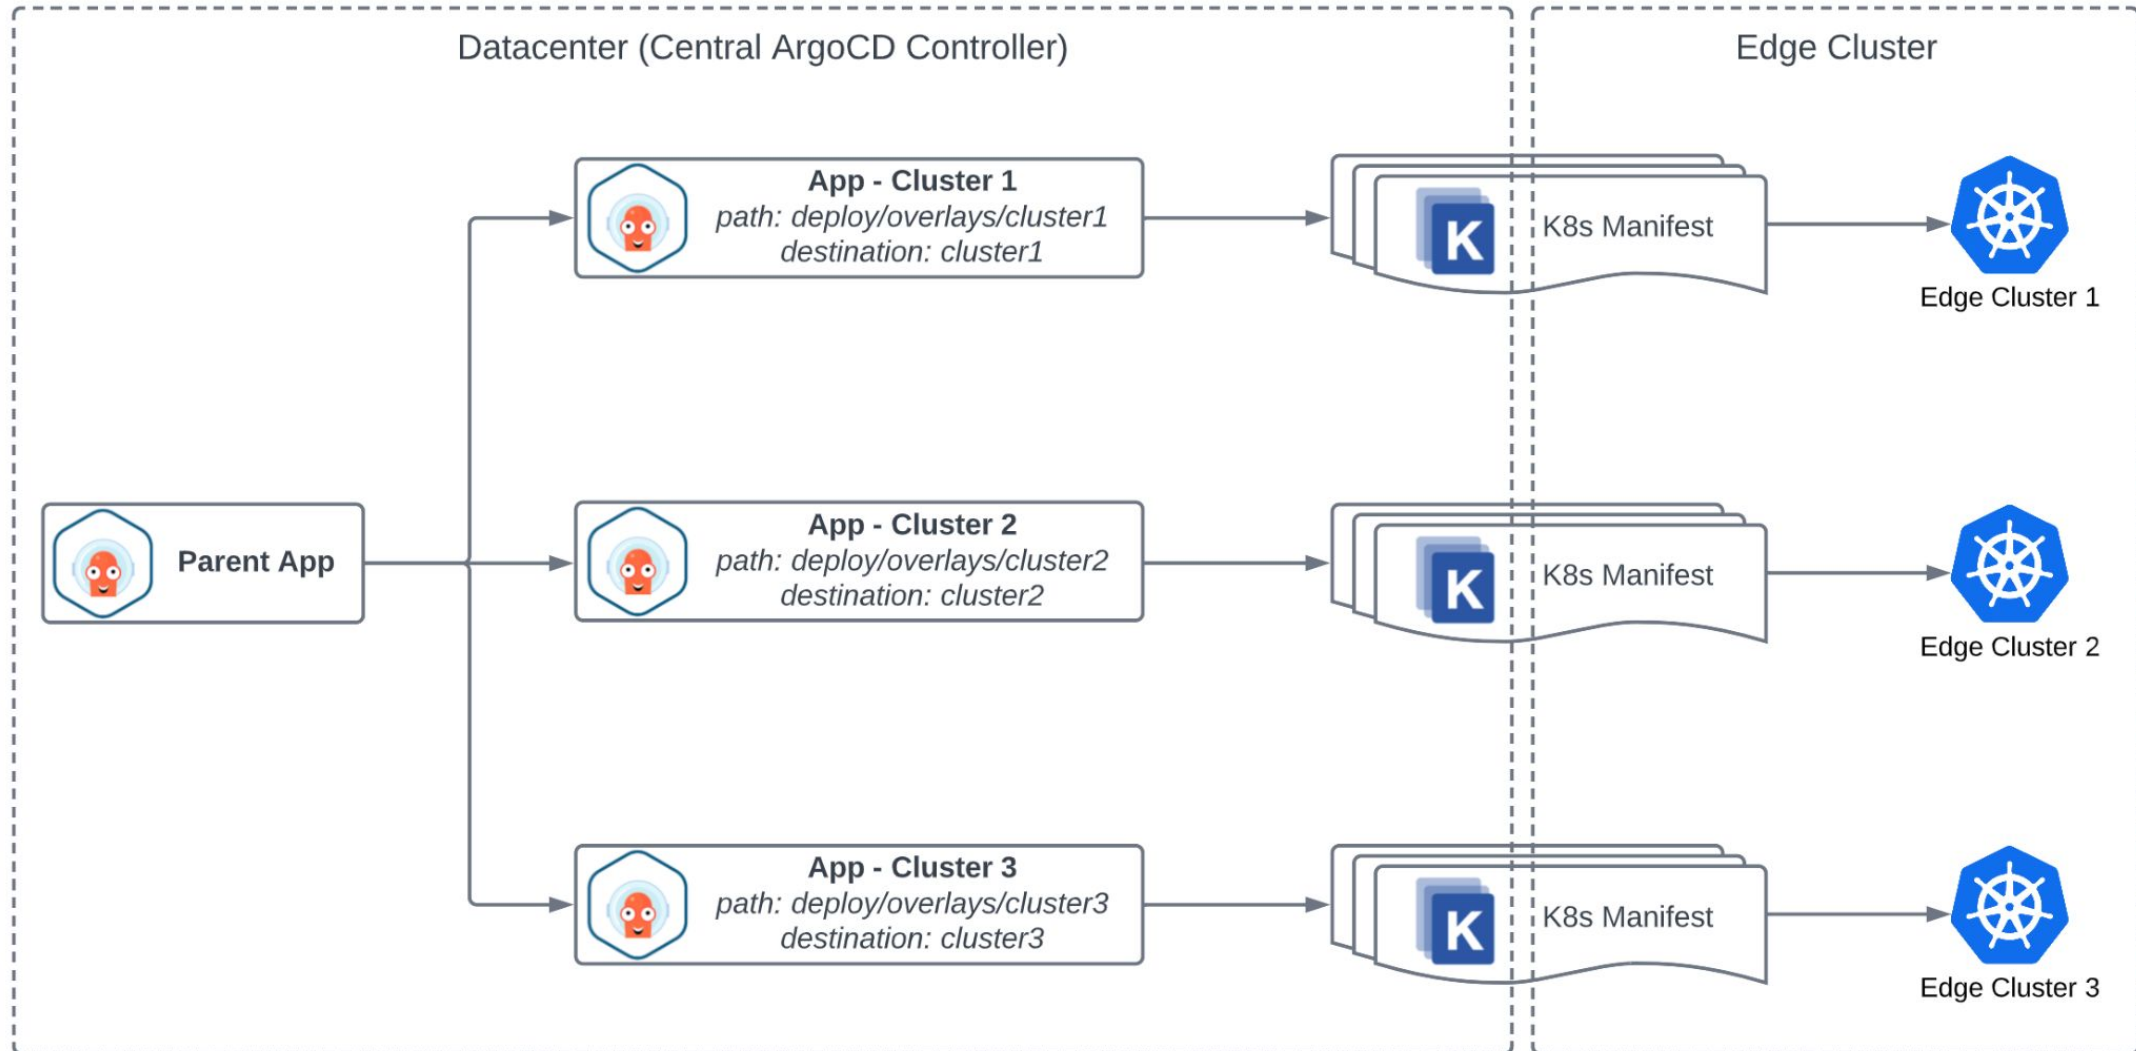
```
1 apiVersion: apps/v1
2 kind: Deployment
3 metadata:
4   name: nginx
5 spec:
6   replicas: 2
7   template:
8     spec:
9       containers:
10      - image: nginx:1.20.0
11        name: nginx
```

```
1 apiVersion: v1
2 kind: Service
3 metadata:
4   name: nginx
5 spec:
6   externalIPs:
7     - 3.3.3.3
```

Edge Deployment Management - Argo CD



- App of apps



Edge Deployment Management - Argo CD




KubeCon



CloudNativeCon

Europe 2022



repo: <https://github.com/qingkunl/argocd-app-of-apps-demo>

Challenges - Argo CD



- Performance
 - Slow to list applications (>10s) when application number gets large (>3000)
- Scalability
 - Imbalanced application handling for different edge clusters
 - Argo CD shards applications to the controller instance by the destination cluster
 - The resource number and type for different edge clusters may vary (big vs small cluster)
- Functionality
 - Need better project-level support (as we use Argo CD project to separate teams/tenants)
 - Application name in different projects could conflict
 - Better project-level view and management, e.g. application, K8s resource, repo, etc.
- Developer Observability
 - Lack of internal observability for tracing, e.g. opentracing
 - Hard to trace and track a single request in source code when troubleshooting

Edge Deployment Models



KubeCon




CloudNativeCon

Europe 2022

Centralized Push

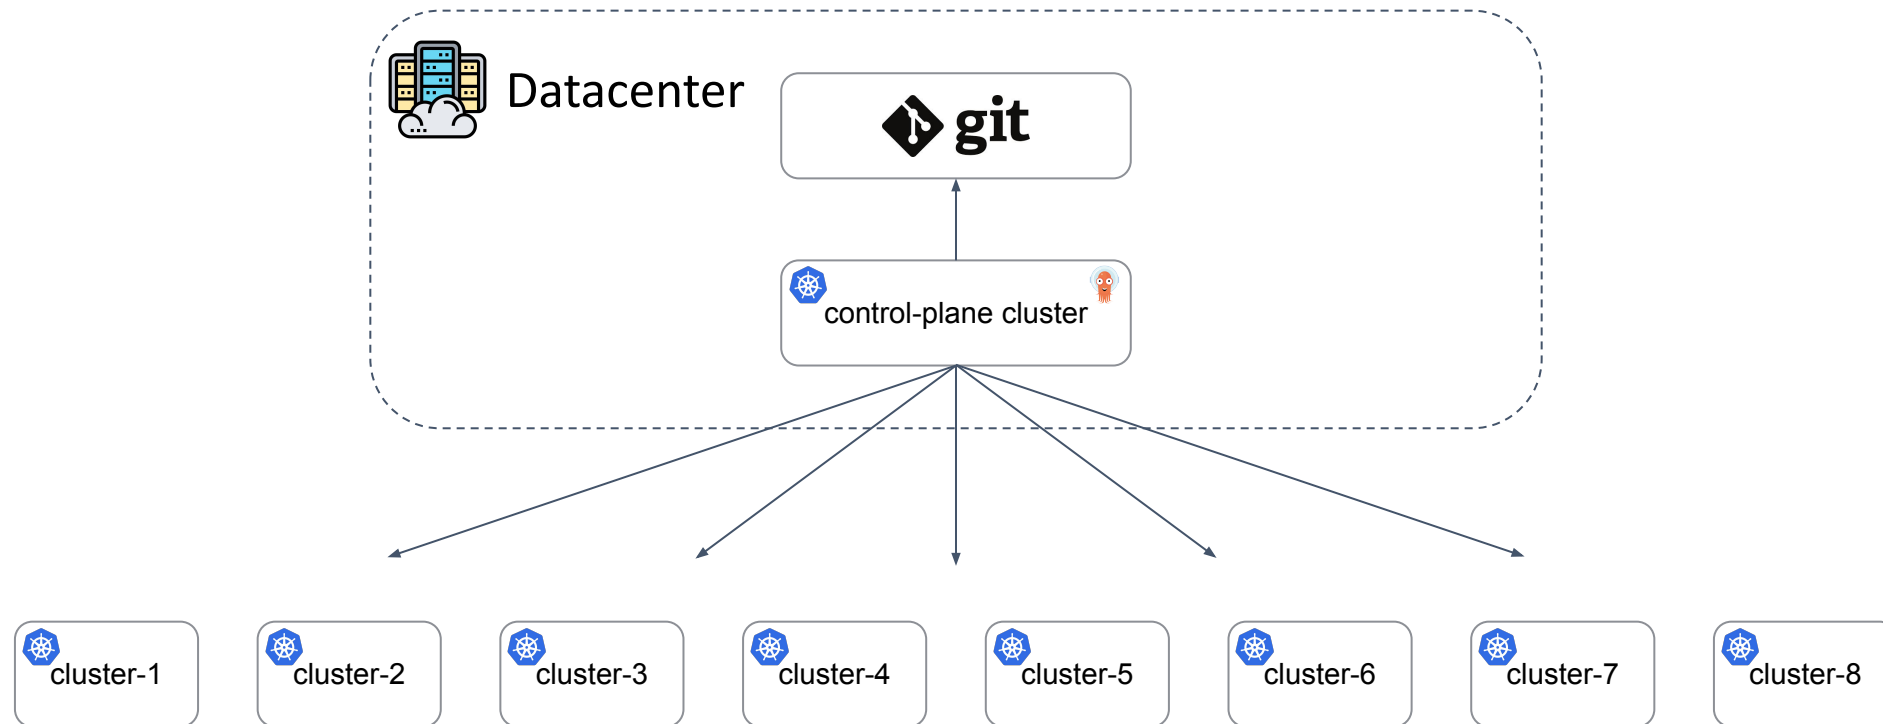
vs.

Distributed Pull



Edge Deployment Model: Centralized Argo CD

Centralized Argo CD *pushes* manifests to individual clusters



Edge Deployment Model: Centralized Argo CD

Advantages

- Single pane of glass - view and control applications across all clusters
- Easiest to manage - single instance to maintain
- API & CLI integrations - automate against Argo CD instance running in datacenter

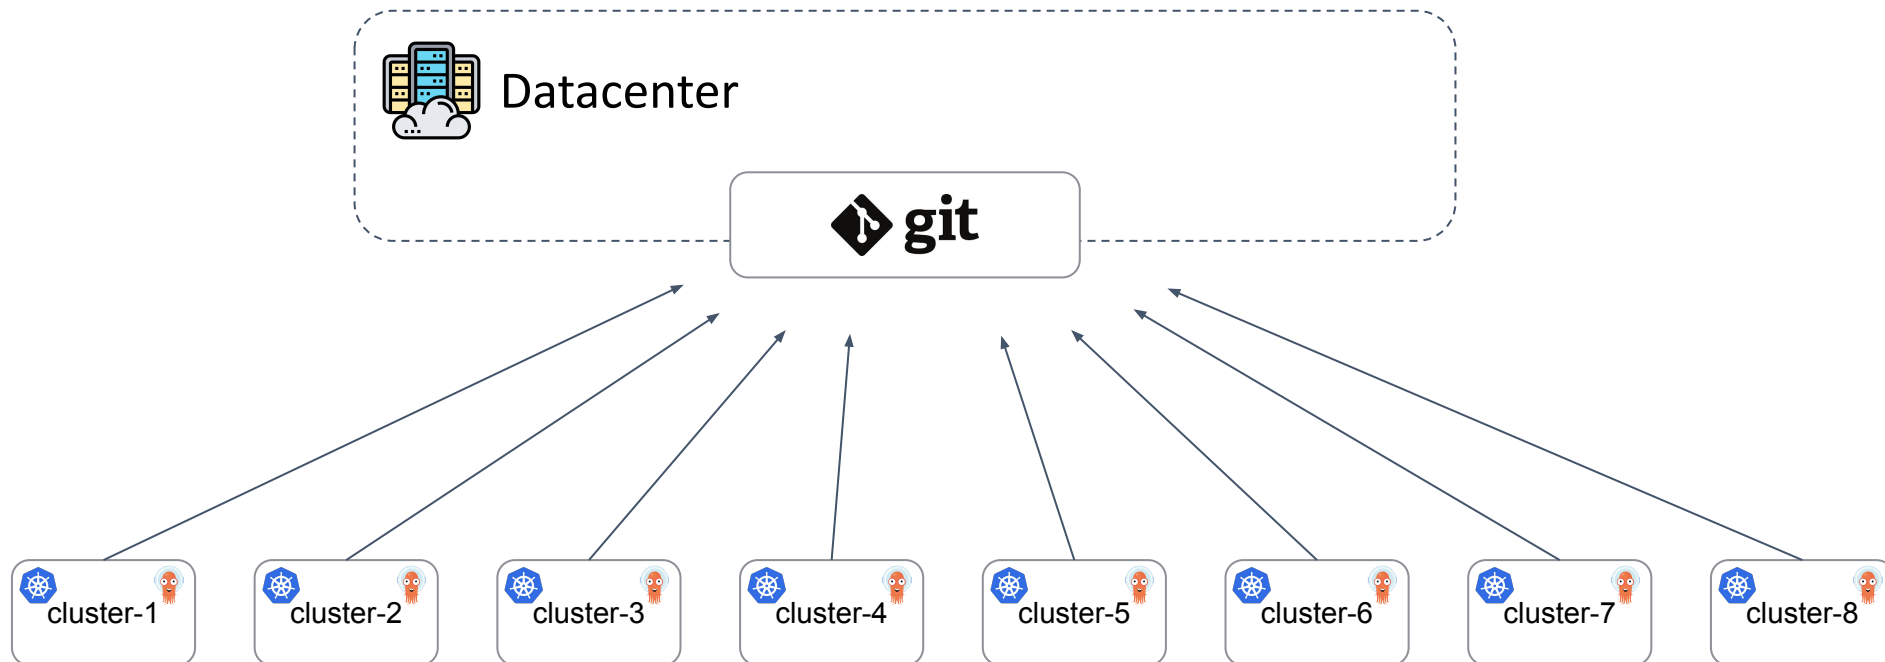
Disadvantages

- Less scalable - requires performance tuning and sharding as cluster fleet grows
- Edge clusters must allow external API Server access
- Single point of failure

Edge Deployment Model: Distributed Argo CD

Argo CD running in edge clusters, ***pulls*** manifests from central git repository

Applications auto-sync changes from cluster specific directories



Edge Deployment Model: Distributed Argo CD

Advantages

- Most scalable - application controller work is distributed to each edge cluster
- Increased security - K8s API servers can be made private

Disadvantages

- No control plane - lack of centralized visibility and management capabilities
- Difficult to manage - many Argo CDs to access/configure/upgrade
- Less flexible - auto-sync requirement reduces flexibility in deploy options
- Git access - clusters requires network access and authentication to the git repository

Edge Techniques: Argo CD Core

A minimal installation mode of Argo CD

- Installs just the minimum necessary components:
 - controller, repo-server, redis
- Great fit for a distributed Argo CD model when UI, API, multi-tenancy are not needed
- UI and CLI are still available with kubeconfig access:

```
export KUBECONFIG=~/.kube/edge-cluster-1
```

```
argocd login --core
```

```
argocd app list
```

```
argocd admin dashboard # available at http://localhost:8080
```

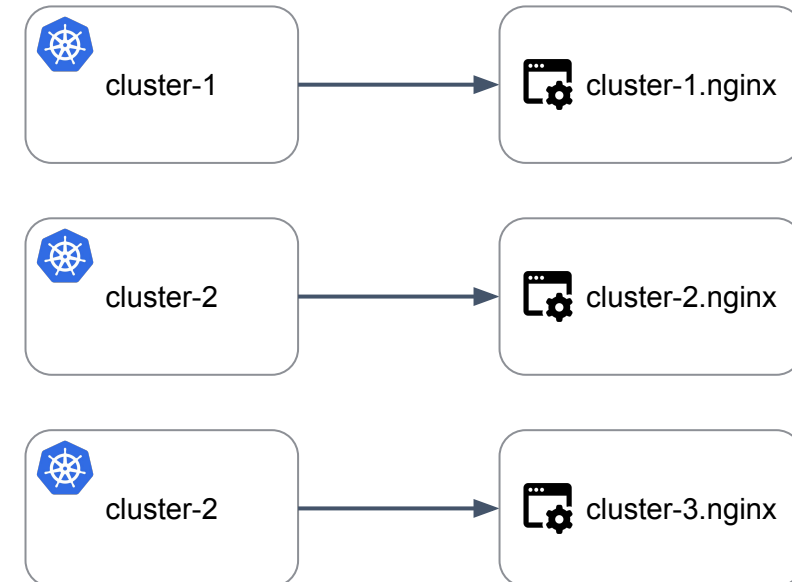
Edge Techniques: ApplicationSets

- Automatically create and deploy applications based on discovery:
 - Paths/files in a git repository
 - Clusters registered to Argo CD
 - Combination of the two
- Created as an alternative to the App-of-Apps
- Works with either the centralized or distributed deployment models

ApplicationSets - Centralized Argo CD

In the central Argo CD, for each cluster registered to Argo CD, create Application(s) for it

```
apiVersion: argoproj.io/v1alpha1
kind: ApplicationSet
metadata:
  name: nginx
spec:
  generators:
  - clusters: {}
  template:
    metadata:
      name: '{{name}}.nginx' # name of cluster
    spec:
      project: "default"
      source:
        repoURL: https://github.com/example/repo
        targetRevision: HEAD
        path: deploy/overlays/{{name}}
      destination:
        server: '{{server}}' # K8s server URL
        namespace: nginx
```



ApplicationSets - Distributed Argo CD

In an edge Argo CD, for each sub directory in a cluster specific path, create Application(s) for it

```
apiVersion: argoproj.io/v1alpha1
kind: ApplicationSet
metadata:
  name: cluster-addons
spec:
  generators:
  - git:
      repoURL: https://github.com/example/repo
      revision: HEAD
      directories:
      - path: deploy/overlays/cluster-3/*
  template:
    metadata:
      name: '{{path.basename}}'
    spec:
      project: default
      source:
        repoURL: https://github.com/example/repo
        targetRevision: HEAD
        path: '{{path}}'
      destination:
        server: https://kubernetes.default.svc
```

deploy


└─ overlays

└─ cluster-3

└─ guestbook

└─ minecraft

└─ nginx



Recent & Future Improvements

Recent Improvements

- Performance and caching improvements
- Tuning options for unreliable networks

Future Improvements

- Applications in different namespaces
 - Prevent application name collisions
 - Better UI performance
- Utility for automatic balancing of cluster shard
- Open Telemetry
- Web shell (kubectl exec in Argo CD UI)



KubeCon



CloudNativeCon

Europe 2022

Thank You!